

Lab-0: Buffer Lab

R Abhijit Srivathsan - 2448044

Compute the dot product of two vectors

```
In [2]: def dot_product(v1, v2):  
        return sum(x * y for x, y in zip(v1, v2))  
  
        # Example  
        v1 = [1, 2, 3]  
        v2 = [4, 5, 6]  
        print("Dot Product:", dot_product(v1, v2))
```

Dot Product: 32

NumPy array (3x3): Transpose, Determinant, Inverse

```
In [3]: import numpy as np  
  
        A = np.array([[1, 2, 3],  
                       [0, 1, 4],  
                       [5, 6, 0]])  
  
        print("Transpose:\n", A.T)  
        print("Determinant:", np.linalg.det(A))  
        print("Inverse:\n", np.linalg.inv(A))
```

Transpose:

```
[[1 0 5]  
 [2 1 6]  
 [3 4 0]]
```

Determinant: 0.99999999999999987

Inverse:

```
[[-24.  18.   5.]  
 [ 20. -15.  -4.]  
 [ -5.   4.   1.]]
```

Load Iris dataset and show statistics

```
In [4]: from sklearn.datasets import load_iris
import pandas as pd

iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)

print(df.describe())
```

| | sepal length (cm) | sepal width (cm) | petal length (cm) | \ |
|-------|-------------------|------------------|-------------------|---|
| count | 150.000000 | 150.000000 | 150.000000 | |
| mean | 5.843333 | 3.057333 | 3.758000 | |
| std | 0.828066 | 0.435866 | 1.765298 | |
| min | 4.300000 | 2.000000 | 1.000000 | |
| 25% | 5.100000 | 2.800000 | 1.600000 | |
| 50% | 5.800000 | 3.000000 | 4.350000 | |
| 75% | 6.400000 | 3.300000 | 5.100000 | |
| max | 7.900000 | 4.400000 | 6.900000 | |


| | petal width (cm) |
|-------|------------------|
| count | 150.000000 |
| mean | 1.199333 |
| std | 0.762238 |
| min | 0.100000 |
| 25% | 0.300000 |
| 50% | 1.300000 |
| 75% | 1.800000 |
| max | 2.500000 |

Scatter plot between two features

```
In [5]: import matplotlib.pyplot as plt

x = df[iris.feature_names[0]]
y = df[iris.feature_names[1]]
labels = iris.target
```

```
plt.figure(figsize=(8, 6))
for i in range(3):
    plt.scatter(x[labels == i], y[labels == i], label=iris.target_names[i])
plt.xlabel(iris.feature_names[0])
plt.ylabel(iris.feature_names[1])
plt.title("Scatter Plot")
plt.legend()
plt.show()
```



Simulate binary classification and compute metrics

```
In [7]: from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score
from sklearn.model_selection import train_test_split
```

```

# Fixing the parameters explicitly
X, y = make_classification(n_samples=100,
                          n_features=2,
                          n_informative=2,
                          n_redundant=0,
                          n_repeated=0,
                          n_classes=2,
                          random_state=42)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))

```

Confusion Matrix:

```
[[14  0]
 [ 0 16]]
```

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

2. Predict student exam performance

Creating a dataset

```

In [8]: import pandas as pd

data = {
    "Study_Hours": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    "Attendance": [60, 65, 70, 75, 80, 85, 90, 95, 95, 100],
    "Sleep_Hours": [6, 6.5, 6, 7, 7.5, 8, 7.5, 6, 5.5, 5],
    "Score": [40, 45, 50, 55, 60, 70, 75, 85, 90, 95]
}

```


```
}  
df = pd.DataFrame(data)
```

a) Load, describe, plot

```
In [9]: print(df.head())  
        print(df.describe())  
  
        plt.figure(figsize=(8, 6))  
        plt.scatter(df["Study_Hours"], df["Score"], color='purple')  
        plt.xlabel("Study Hours")  
        plt.ylabel("Score")  
        plt.title("Study Hours vs Score")  
        plt.grid(True)  
        plt.show()
```

| | Study_Hours | Attendance | Sleep_Hours | Score |
|---|-------------|------------|-------------|-------|
| 0 | 1 | 60 | 6.0 | 40 |
| 1 | 2 | 65 | 6.5 | 45 |
| 2 | 3 | 70 | 6.0 | 50 |
| 3 | 4 | 75 | 7.0 | 55 |
| 4 | 5 | 80 | 7.5 | 60 |

| | Study_Hours | Attendance | Sleep_Hours | Score |
|-------|-------------|------------|-------------|-----------|
| count | 10.00000 | 10.000000 | 10.000000 | 10.000000 |
| mean | 5.50000 | 81.500000 | 6.500000 | 66.500000 |
| std | 3.02765 | 13.753787 | 0.971825 | 19.443651 |
| min | 1.00000 | 60.000000 | 5.000000 | 40.000000 |
| 25% | 3.25000 | 71.250000 | 6.000000 | 51.250000 |
| 50% | 5.50000 | 82.500000 | 6.250000 | 65.000000 |
| 75% | 7.75000 | 93.750000 | 7.375000 | 82.500000 |
| max | 10.00000 | 100.000000 | 8.000000 | 95.000000 |



b) Matrix form: mean, covariance, eigenvalues

```
In [10]: X = df[["Study_Hours", "Attendance", "Sleep_Hours", "Score"]].values  
  
mean_vector = np.mean(X, axis=0)  
cov_matrix = np.cov(X.T)
```

```
eigenvalues, _ = np.linalg.eig(cov_matrix)

print("Mean Vector:", mean_vector)
print("Covariance Matrix:\n", cov_matrix)
print("Eigenvalues:", eigenvalues)
```

Mean Vector: [5.5 81.5 6.5 66.5]

Covariance Matrix:

```
[[ 9.16666667 41.38888889 -0.77777778 58.61111111]
 [41.38888889 189.16666667 -2.5          264.16666667]
 [-0.77777778 -2.5          0.94444444 -5.83333333]
 [58.61111111 264.16666667 -5.83333333 378.05555556]]
```

Eigenvalues: [5.73347811e+02 3.68705167e+00 3.62017829e-02 2.62269129e-01]


c) Histogram, boxplot, and descriptive stats

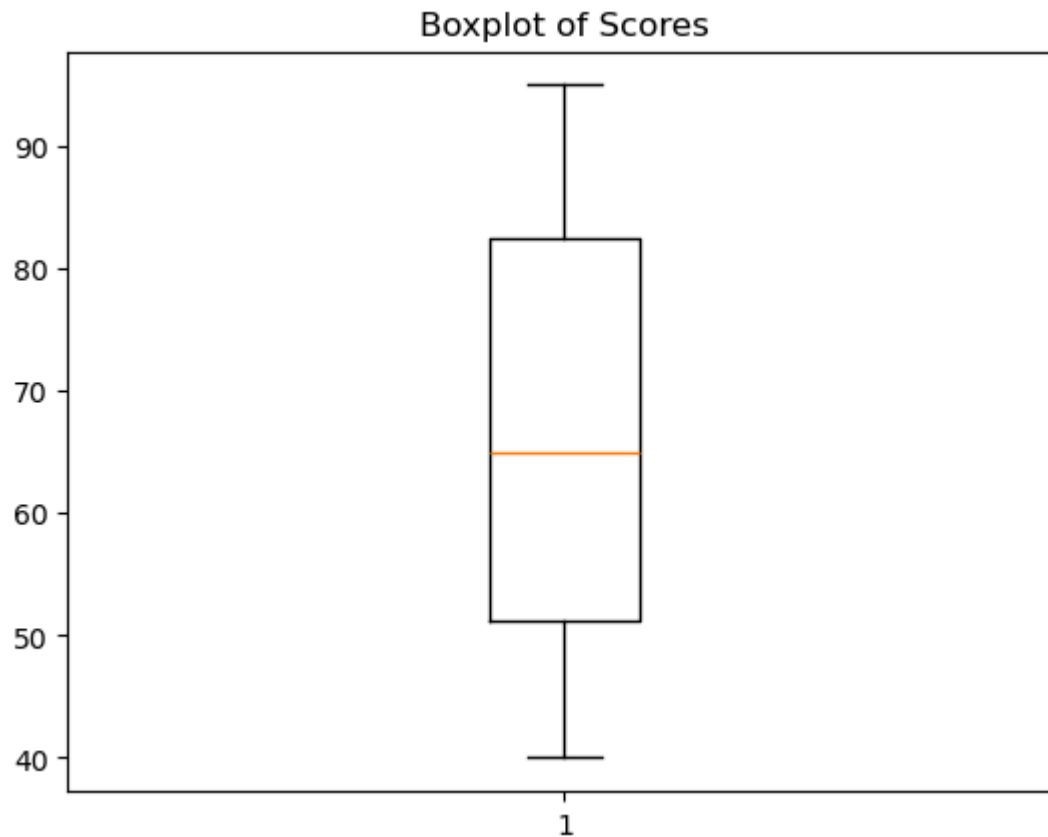
```
In [11]: from scipy.stats import skew, kurtosis, mode

plt.hist(df["Score"], bins=5, color="skyblue")
plt.title("Histogram of Scores")
plt.show()

plt.boxplot(df["Score"])
plt.title("Boxplot of Scores")
plt.show()

print("Mean:", df["Score"].mean())
print("Median:", df["Score"].median())
print("Mode:", mode(df["Score"], keepdims=False).mode)
print("Standard Deviation:", df["Score"].std())
print("Skewness:", skew(df["Score"]))
print("Kurtosis:", kurtosis(df["Score"]))
```



Mean: 66.5
Median: 65.0
Mode: 40
Standard Deviation: 19.443650777453175
Skewness: 0.13192683268558208
Kurtosis: -1.369647053615437

d) Train/test split, Linear Regression, metrics, plot

```
In [12]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split

X = df[["Study_Hours", "Attendance", "Sleep_Hours"]]
```

```
y = df["Score"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))

# Plot actual vs predicted
plt.scatter(y_test, y_pred, color='green')
plt.xlabel("Actual Scores")
plt.ylabel("Predicted Scores")
plt.title("Actual vs Predicted Scores")
plt.plot([y.min(), y.max()], [y.min(), y.max()], color='red', linestyle='--')
plt.grid(True)
plt.show()
```

Mean Squared Error: 3.6939959685219503

R² Score: 0.9927032178399566

Actual vs Predicted Scores

