Lab-4: Single-Layer Neural Network Alarm System

R Abhijit Srivathsan - 2448044

```
In [1]: import numpy as np
    np.set_printoptions(precision=4, suppress=True)
    np.random.seed(42)
```

Mathematical Foundation

This implementation demonstrates the fundamental concepts of neural network training:

Network Architecture

- Input Layer: 2 neurons (x₁, x₂) representing sensor inputs
- Output Layer: 1 neuron with sigmoid activation producing alarm signal
- Parameters: 2 weights (w_1, w_2) and 1 bias (b)

Forward Propagation

The network computes:

- 1. Linear combination: $z = w_1x_1 + w_2x_2 + b$
- 2. **Activation**: $\hat{y} = \sigma(z) = 1/(1 + e^{-z})$

Loss Function

Mean Squared Error (MSE): $E = (1/n)\Sigma(y - \hat{y})^2$

Learning Algorithm

Delta Rule with batch gradient descent:

```
Weight update: w ← w - η(∂E/∂w)
Bias update: b ← b - η(∂E/∂b)
Learning rate: η = 0.1
```

1. Training data

For this scenario we assume that **any** sensor being active should raise the alarm (logical **OR**).

Truth Table for Alarm System

The training data represents all possible combinations of sensor inputs with their expected alarm outputs:

x ₁ (Motion)	X₂ (Door)	y (Alarm)	Logic
0	0	0	No sensors active → No alarm
0	1	1	Door open → Alarm
1	0	1	Motion detected → Alarm
1	1	1	Both sensors → Alarm

This implements **logical OR** behavior: the alarm should sound if **any** sensor detects activity. This is a linearly separable problem that a single-layer perceptron can solve.

2. Helper functions

```
In [3]: def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def sigmoid_derivative(a):
    return a * (1 - a)
```

Sigmoid Activation Function

The **sigmoid function** $\sigma(z) = 1/(1 + e^{-z})$ has several important properties:

- Range: (0, 1) perfect for binary classification
- **Differentiable**: Smooth gradient for backpropagation
- S-shaped curve: Provides non-linear decision boundary
- **Derivative**: $\sigma'(z) = \sigma(z)(1 \sigma(z))$ computationally efficient

The sigmoid derivative is used during backpropagation to compute gradients for weight updates.

3. Initialize weights and bias (small random values)

```
In [4]: # two inputs -> one output
    w = np.random.randn(2, 1) * 0.1 # (2x1)
    b = np.random.randn(1) * 0.1 # scalar bias
    print("Initial weights:\n", w)
    print("Initial bias:\n", b)

Initial weights:
    [[ 0.0497]
    [-0.0138]]
    Initial bias:
    [0.0648]
```



Weight Initialization Strategy

Small Random Initialization is used here with the following rationale:

- Random values: Break symmetry if all weights start identical, they remain identical
- Small magnitude (×0.1): Prevents saturation of sigmoid function
- Normal distribution: Provides balanced positive/negative initial values
- Reproducible: np.random.seed(42) ensures consistent results

The initial weights and bias will be small random values around zero, allowing the network to learn from gradient descent updates.

4. Forward pass – predictions and intermediate values

```
In [5]: # Linear combination
    z = np.dot(data_x, w) + b # (4x1)
    # Activation
    y_hat = sigmoid(z)
    print("z (linear output):\n", z)
    print("ŷ (sigmoid output):\n", y_hat)

# Error (Mean Squared Error)
    error = np.mean((data_y - y_hat) ** 2)
    print("\nMean Squared Error before update =", error)
```

```
z (linear output):
[[0.0648]
[0.0509]
 [0.1144]
[0.1006]]
ŷ (sigmoid output):
[[0.5162]
[0.5127]
 [0.5286]
[0.5251]]
Mean Squared Error before update = 0.23790376766405488
 Forward Pass Output & Error (Before Update)
   • Linear Combination (z): [ z = X \cdot W + b =
                                                             0.0648
                                                             0.0509
                                                             0.1144
                                                             [0.1006]
   • Sigmoid Activation (ŷ): [\hat{y} = \sigma(z) =
                                                             0.5162
                                                             0.5127
                                                             0.5286
                                                             \lfloor\,0.5251\,
floor
   • Mean Squared Error (before update):
     [\text{text}\{MSE\} = 0.2379]
```

Forward Pass Analysis

The forward pass computes predictions for all training samples simultaneously:

- 1. Linear Transformation: z = Xw + b
 - Matrix multiplication: $(4\times2)\times(2\times1)+(1\times1)=(4\times1)$
 - Each row represents one training sample's linear output
- 2. **Activation**: $\hat{y} = \sigma(z)$
 - Applies sigmoid element-wise to convert linear outputs to probabilities
 - Values close to 0 indicate "no alarm", close to 1 indicate "alarm"
- 3. Error Calculation: MSE = mean($(y \hat{y})^2$)
 - Measures how far predictions are from true labels
 - Higher MSE indicates poorer performance

The initial predictions will likely be poor since weights are random, but this establishes our baseline before learning.

5. Backward pass – update with delta rule

```
In [6]: # Compute gradient of error w.r.t. y_hat
    error_grad = -(data_y - y_hat)  # d(MSE)/dŷ for each sample (4x1)
# Gradient w.r.t. z
z_grad = error_grad * sigmoid_derivative(y_hat)

# Gradients for weights and bias (batch = mean of sample grads)
w_grad = np.dot(data_x.T, z_grad) / len(data_x)
b_grad = np.mean(z_grad)

lr = 0.1
w_new = w - lr * w_grad
b_new = b - lr * b_grad

print("Gradient w.r.t weights:\n", w_grad)
```

```
print("Gradient w.r.t bias:\n", b grad)
 print("\nUpdated weights:\n", w new)
 print("Updated bias:\n", b new)
Gradient w.r.t weights:
[[-0.059]
[-0.06]]
Gradient w.r.t bias:
 -0.05717840777138376
Updated weights:
[[ 0.0556]
[-0.0078]]
Updated bias:
 [0.0705]
 Backpropagation & Weight Update
   • Gradient w.r.t. Weights: [\frac{\partial \text{MSE}}{\partial W} =
   • Gradient w.r.t. Bias: [\frac{\partial \text{MSE}}{\partial b} = -0.0572 ]
   • Updated Weights (after learning rate = 0.1): [ W_{\text{new}} =
                                                             0.0556
                                                            -0.0078
   • Updated Bias: [b {\text{new}} = 0.0705]
```

Backpropagation and Weight Update

The **Delta Rule** implements gradient descent to minimize error:

Step 1: Error Gradient

- error_grad = $-(y \hat{y})$ Derivative of MSE w.r.t. predictions
- Negative sign because we want to minimize error

Step 2: Local Gradient

- $z_grad = error_grad \times \sigma'(\hat{y})$ Chain rule application
- Sigmoid derivative weights the error by activation sensitivity

Step 3: Parameter Gradients

- Weight gradient: $\partial E/\partial w = X^T \times z \text{ grad } / \text{ n}$ (averaged over batch)
- Bias gradient: $\partial E/\partial b = mean(z grad)$ (averaged over batch)

Step 4: Parameter Update

- New weights: $W \leftarrow W \eta \times \partial E/\partial W$
- New bias: $b \leftarrow b \eta \times \partial E/\partial b$
- Learning rate $\eta = 0.1$ controls step size

This single update should move the network toward better OR gate behavior.

6. Verifying new error after one update (optional)

```
In [7]: # Forward pass with updated parameters
z2 = np.dot(data_x, w_new) + b_new
y_hat2 = sigmoid(z2)
error2 = np.mean((data_y - y_hat2) ** 2)
print("New Mean Squared Error =", error2)
```

New Mean Squared Error = 0.23584379921143195

▼ Forward Pass After Update

- New Mean Squared Error: [\text{MSE}_{\text{after update}}} = 0.2358]
 - ▼ The error decreased from **0.2379** to **0.2358**, showing that the network has started learning and adjusting weights in the right direction.

Learning Progress Analysis

After one epoch of training, we can observe:

- 1. Error Reduction: The new MSE should be lower than the initial MSE, indicating learning
- 2. Weight Evolution: Updated weights should better reflect the OR logic pattern
- 3. Prediction Improvement: New predictions should be closer to desired outputs

Expected Behavior:

- For input [0,0]: prediction should move toward 0 (no alarm)
- For inputs [0,1], [1,0], [1,1]: predictions should move toward 1 (alarm)

Convergence Notes:

- Single-layer networks can perfectly learn linearly separable patterns like OR
- Multiple epochs would be needed for complete convergence
- The delta rule guarantees convergence for linearly separable problems

🧠 Final Inference & Conclusion

After performing one forward and backward pass on a single-layer neural network with a sigmoid activation:

- The model learned to reduce its prediction error by adjusting weights and bias.
- A clear reduction in Mean Squared Error (MSE) confirms that backpropagation worked as expected.
- This demonstrates the basic principle of learning in neural networks using gradient descent and the delta rule.

Error Comparison Table

Step	MSE	Weights	Bias
Before Update	0.2379	[[0.0497], [-0.0138]]	0.0648
After One Update Step	0.2358	[[0.0556], [-0.0078]]	0.0705

Conclusion:

The decrease in MSE indicates successful learning. This basic experiment confirms that a single-layer neural network can learn to map inputs to outputs by minimizing error using gradient descent and backpropagation.