

Lab-5: CIFAR-10 Image Classification: Custom CNN vs AlexNet

R Abhijit Srivathsan - 2448044

Project Overview

This notebook implements and compares two approaches for CIFAR-10 image classification:

1. Custom CNN architecture
2. AlexNet architecture adapted for CIFAR-10

Dataset: CIFAR-10

- 60,000 32x32 color images in 10 classes
- 50,000 training images and 10,000 test images
- Classes: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck

```
In [1]: # Import required libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, models
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix
import time

# Set random seeds for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

print(f"TensorFlow version: {tf.__version__}")
print(f"GPU Available: {tf.config.list_physical_devices('GPU')}")
```

```

2025-07-28 20:13:08.597217: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-07-28 20:13:08.606302: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1753713788.616927 16550 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1753713788.620016 16550 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
W0000 00:00:1753713788.627938 16550 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1753713788.627954 16550 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1753713788.627956 16550 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1753713788.627957 16550 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
2025-07-28 20:13:08.630928: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
/home/abhijit/miniconda3/envs/tf-env/lib/python3.12/site-packages/requests/__init__.py:86: RequestsDependencyWarning: Unable to find acceptable character detection dependency (chardet or charset_normalizer).
  warnings.warn(
TensorFlow version: 2.19.0
GPU Available: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]

```

```

In [2]: # Load and preprocess CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()

# Class names for CIFAR-10
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

print(f"Training data shape: {x_train.shape}")
print(f"Training labels shape: {y_train.shape}")
print(f"Test data shape: {x_test.shape}")
print(f"Test labels shape: {y_test.shape}")
print(f"Number of classes: {len(class_names)}")

```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

170498071/170498071  **529s** 3us/step

Training data shape: (50000, 32, 32, 3)

Training labels shape: (50000, 1)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000, 1)

Number of classes: 10

```
In [3]: # Visualize sample images from CIFAR-10
plt.figure(figsize=(12, 8))
for i in range(20):
    plt.subplot(4, 5, i + 1)
    plt.imshow(x_train[i])
    plt.title(f'{class_names[y_train[i][0]]}')
    plt.axis('off')
plt.suptitle('Sample CIFAR-10 Images', fontsize=16)
plt.tight_layout()
plt.show()
```

Sample CIFAR-10 Images

frog



truck



truck



deer



automobile



automobile



bird



horse



ship



cat



deer




horse



horse




bird



truck



truck




truck




cat



bird



frog



```
In [4]: # Data preprocessing
# Normalize pixel values to [0, 1]
x_train_normalized = x_train.astype('float32') / 255.0
x_test_normalized = x_test.astype('float32') / 255.0

# Convert labels to categorical (one-hot encoding)
y_train_categorical = keras.utils.to_categorical(y_train, 10)
y_test_categorical = keras.utils.to_categorical(y_test, 10)

print(f"Normalized training data range: [{x_train_normalized.min():.2f}, {x_train_normalized.max():.2f}]")
print(f"Label shape after one-hot encoding: {y_train_categorical.shape}")
```

Normalized training data range: [0.00, 1.00]
Label shape after one-hot encoding: (50000, 10)

Model 1: Custom CNN Architecture

```
In [5]: def create_custom_cnn():
        """
        Create a custom CNN architecture optimized for CIFAR-10
        """
        model = models.Sequential([
            # First Convolutional Block
            layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3), padding='same'),
            layers.BatchNormalization(),
            layers.Conv2D(32, (3, 3), activation='relu', padding='same'),
            layers.MaxPooling2D((2, 2)),
            layers.Dropout(0.25),

            # Second Convolutional Block
            layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
            layers.BatchNormalization(),
            layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
            layers.MaxPooling2D((2, 2)),
            layers.Dropout(0.25),

            # Third Convolutional Block
            layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
            layers.BatchNormalization(),
```

```

layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
layers.MaxPooling2D((2, 2)),
layers.Dropout(0.25),

# Fully Connected Layers
layers.Flatten(),
layers.Dense(512, activation='relu'),
layers.BatchNormalization(),
layers.Dropout(0.5),
layers.Dense(256, activation='relu'),
layers.Dropout(0.5),
layers.Dense(10, activation='softmax')
])

return model

# Create and compile custom CNN
custom_cnn = create_custom_cnn()
custom_cnn.compile(optimizer='adam',
                   loss='categorical_crossentropy',
                   metrics=['accuracy'])

print("Custom CNN Architecture:")
custom_cnn.summary()

```

/home/abhijit/miniconda3/envs/tf-env/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)
I0000 00:00:1753714324.161874 16550 gpu_device.cc:2019] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 3775 MB memory: -> device: 0, name: NVIDIA GeForce RTX 4050 Laptop GPU, pci bus id: 0000:01:00.0, compute capability: 8.9

Custom CNN Architecture:

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (BatchNormalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9,248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36,928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 8, 8, 128)	512
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147,584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_2 (Dropout)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1,049,088
batch_normalization_3	(None, 512)	2,048

(BatchNormalization)		
dropout_3 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328
dropout_4 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2,570

Total params: 1,472,938 (5.62 MB)

Trainable params: 1,471,466 (5.61 MB)

Non-trainable params: 1,472 (5.75 KB)

Model 2: AlexNet Architecture (Adapted for CIFAR-10)

```
In [6]: def create_alexnet_cifar10():
        """
        Create AlexNet architecture adapted for CIFAR-10 (32x32 input)
        Original AlexNet was designed for 224x224 images, so we adapt it for smaller CIFAR-10 images
        """
        model = models.Sequential([
            # First Convolutional Layer
            layers.Conv2D(96, (5, 5), strides=(1, 1), activation='relu', input_shape=(32, 32, 3), padding='same'),
            layers.BatchNormalization(),
            layers.MaxPooling2D((2, 2), strides=(2, 2)),

            # Second Convolutional Layer
            layers.Conv2D(256, (5, 5), activation='relu', padding='same'),
            layers.BatchNormalization(),
            layers.MaxPooling2D((2, 2), strides=(2, 2)),

            # Third Convolutional Layer
            layers.Conv2D(384, (3, 3), activation='relu', padding='same'),

            # Fourth Convolutional Layer
            layers.Conv2D(384, (3, 3), activation='relu', padding='same'),
```



```

        # Fifth Convolutional Layer
        layers.Conv2D(256, (3, 3), activation='relu', padding='same'),
        layers.MaxPooling2D((2, 2), strides=(2, 2)),

        # Fully Connected Layers
        layers.Flatten(),
        layers.Dense(4096, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(4096, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(10, activation='softmax')
    ])

    return model

# Create and compile AlexNet
alexnet_model = create_alexnet_cifar10()
alexnet_model.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])

print("AlexNet Architecture (Adapted for CIFAR-10):")
alexnet_model.summary()

```

AlexNet Architecture (Adapted for CIFAR-10):
Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 96)	7,296
batch_normalization_4 (BatchNormalization)	(None, 32, 32, 96)	384
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 96)	0
conv2d_7 (Conv2D)	(None, 16, 16, 256)	614,656
batch_normalization_5 (BatchNormalization)	(None, 16, 16, 256)	1,024
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 256)	0
conv2d_8 (Conv2D)	(None, 8, 8, 384)	885,120
conv2d_9 (Conv2D)	(None, 8, 8, 384)	1,327,488
conv2d_10 (Conv2D)	(None, 8, 8, 256)	884,992
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 256)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_3 (Dense)	(None, 4096)	16,781,312
dropout_5 (Dropout)	(None, 4096)	0
dense_4 (Dense)	(None, 4096)	16,781,312
dropout_6 (Dropout)	(None, 4096)	0
dense_5 (Dense)	(None, 10)	40,970

Total params: 37,324,554 (142.38 MB)

Trainable params: 37,323,850 (142.38 MB)

Non-trainable params: 704 (2.75 KB)

Training Configuration and Callbacks

```
In [7]: # Training configuration
EPOCHS = 50
BATCH_SIZE = 128
VALIDATION_SPLIT = 0.1

# Create callbacks for better training
def create_callbacks(model_name):
    return [
        keras.callbacks.EarlyStopping(
            monitor='val_accuracy',
            patience=10,
            restore_best_weights=True
        ),
        keras.callbacks.ReduceLROnPlateau(
            monitor='val_loss',
            factor=0.2,
            patience=5,
            min_lr=0.0001
        ),
        keras.callbacks.ModelCheckpoint(
            f'best_{model_name}_model.h5',
            monitor='val_accuracy',
            save_best_only=True,
            verbose=1
        )
    ]
```

Training Custom CNN

```
In [8]: print("Training Custom CNN...")
start_time = time.time()

# Train custom CNN
custom_cnn_history = custom_cnn.fit(
    x_train_normalized, y_train_categorical,
```

```

    batch_size=BATCH_SIZE,
    epochs=EPOCHS,
    validation_split=VALIDATION_SPLIT,
    callbacks=create_callbacks('custom_cnn'),
    verbose=1
)


custom_cnn_training_time = time.time() - start_time
print(f"Custom CNN training completed in {custom_cnn_training_time:.2f} seconds")

```


Training Custom CNN...

Epoch 1/50


WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1753714327.019846 20630 service.cc:152] XLA service 0x76744c007380 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:
I0000 00:00:1753714327.019869 20630 service.cc:160] StreamExecutor device (0): NVIDIA GeForce RTX 4050 Laptop GPU, Compute Capability 8.9
2025-07-28 20:22:07.077048: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
I0000 00:00:1753714327.430286 20630 cuda_dnn.cc:529] Loaded cuDNN version 90300

11/352  **5s** 17ms/step - accuracy: 0.1285 - loss: 3.7682

I0000 00:00:1753714334.588397 20630 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

349/352  **0s** 15ms/step - accuracy: 0.2753 - loss: 2.3329

2025-07-28 20:22:21.261254: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_3803', 60 bytes spill stores, 60 bytes spill loads


352/352  **0s** 32ms/step - accuracy: 0.2759 - loss: 2.3293

2025-07-28 20:22:26.415992: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_207', 4 bytes spill stores, 4 bytes spill loads


2025-07-28 20:22:26.426765: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_207', 4 bytes spill stores, 4 bytes spill loads

Epoch 1: val_accuracy improved from -inf to 0.12880, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  24s 42ms/step - accuracy: 0.2761 - loss: 2.3281 - val_accuracy: 0.1288 - val_loss: 3.6880 - learning_rate: 0.0010

Epoch 2/50


351/352  0s 16ms/step - accuracy: 0.4781 - loss: 1.4330

Epoch 2: val_accuracy improved from 0.12880 to 0.58760, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 17ms/step - accuracy: 0.4783 - loss: 1.4325 - val_accuracy: 0.5876 - val_loss: 1.1455 - learning_rate: 0.0010

Epoch 3/50


352/352  0s 16ms/step - accuracy: 0.5825 - loss: 1.1629

Epoch 3: val_accuracy improved from 0.58760 to 0.65720, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 17ms/step - accuracy: 0.5826 - loss: 1.1628 - val_accuracy: 0.6572 - val_loss: 0.9807 - learning_rate: 0.0010

Epoch 4/50


349/352  0s 17ms/step - accuracy: 0.6506 - loss: 0.9878

Epoch 4: val_accuracy improved from 0.65720 to 0.71900, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  6s 18ms/step - accuracy: 0.6507 - loss: 0.9875 - val_accuracy: 0.7190 - val_loss: 0.7975 - learning_rate: 0.0010


Epoch 5/50

351/352  0s 17ms/step - accuracy: 0.6896 - loss: 0.8752

Epoch 5: val_accuracy did not improve from 0.71900


352/352  6s 17ms/step - accuracy: 0.6896 - loss: 0.8751 - val_accuracy: 0.7190 - val_loss: 0.8153 - learning_rate: 0.0010

Epoch 6/50


352/352  0s 17ms/step - accuracy: 0.7208 - loss: 0.8054

Epoch 6: val_accuracy improved from 0.71900 to 0.74840, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 17ms/step - accuracy: 0.7208 - loss: 0.8053 - val_accuracy: 0.7484 - val_loss: 0.7164 - learning_rate: 0.0010

Epoch 7/50


350/352  0s 17ms/step - accuracy: 0.7437 - loss: 0.7378

Epoch 7: val_accuracy improved from 0.74840 to 0.76560, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 18ms/step - accuracy: 0.7438 - loss: 0.7377 - val_accuracy: 0.7656 - val_loss: 0.7203 - learning_rate: 0.0010

Epoch 8/50


349/352  0s 17ms/step - accuracy: 0.7590 - loss: 0.6984

Epoch 8: val_accuracy improved from 0.76560 to 0.77280, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 18ms/step - accuracy: 0.7590 - loss: 0.6982 - val_accuracy: 0.7728 - val_loss: 0.6670 - learning_rate: 0.0010

Epoch 9/50


351/352  0s 17ms/step - accuracy: 0.7767 - loss: 0.6404

Epoch 9: val_accuracy improved from 0.77280 to 0.77940, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 17ms/step - accuracy: 0.7768 - loss: 0.6403 - val_accuracy: 0.7794 - val_loss: 0.6639 - learning_rate: 0.0010


Epoch 10/50

350/352  0s 17ms/step - accuracy: 0.7897 - loss: 0.6095

Epoch 10: val_accuracy improved from 0.77940 to 0.79800, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  6s 18ms/step - accuracy: 0.7897 - loss: 0.6094 - val_accuracy: 0.7980 - val_loss: 0.6023 - learning_rate: 0.0010
Epoch 11/50

350/352  0s 17ms/step - accuracy: 0.8004 - loss: 0.5823


Epoch 11: val_accuracy improved from 0.79800 to 0.80620, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  6s 18ms/step - accuracy: 0.8005 - loss: 0.5822 - val_accuracy: 0.8062 - val_loss: 0.5710 - learning_rate: 0.0010
Epoch 12/50

349/352  0s 17ms/step - accuracy: 0.8111 - loss: 0.5496


Epoch 12: val_accuracy did not improve from 0.80620


352/352  6s 17ms/step - accuracy: 0.8111 - loss: 0.5495 - val_accuracy: 0.7980 - val_loss: 0.6106 - learning_rate: 0.0010
Epoch 13/50

351/352  0s 17ms/step - accuracy: 0.8224 - loss: 0.5189


Epoch 13: val_accuracy improved from 0.80620 to 0.81600, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  6s 18ms/step - accuracy: 0.8224 - loss: 0.5188 - val_accuracy: 0.8160 - val_loss: 0.5592 - learning_rate: 0.0010
Epoch 14/50


350/352  0s 17ms/step - accuracy: 0.8320 - loss: 0.4929


Epoch 14: val_accuracy did not improve from 0.81600

352/352  6s 17ms/step - accuracy: 0.8320 - loss: 0.4929 - val_accuracy: 0.8114 - val_loss: 0.5848 - learning_rate: 0.0010
Epoch 15/50

350/352  0s 17ms/step - accuracy: 0.8350 - loss: 0.4794


Epoch 15: val_accuracy did not improve from 0.81600


352/352  6s 17ms/step - accuracy: 0.8351 - loss: 0.4793 - val_accuracy: 0.8112 - val_loss: 0.5624 - learning_rate: 0.0010
Epoch 16/50


352/352  0s 17ms/step - accuracy: 0.8424 - loss: 0.4591


Epoch 16: val_accuracy improved from 0.81600 to 0.82940, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **6s** 17ms/step - accuracy: 0.8424 - loss: 0.4591 - val_accuracy: 0.8294 - val_loss: 0.5167 - learning_rate: 0.0010
Epoch 17/50

349/352  **0s** 17ms/step - accuracy: 0.8510 - loss: 0.4357
Epoch 17: val_accuracy did not improve from 0.82940

352/352  **6s** 17ms/step - accuracy: 0.8510 - loss: 0.4356 - val_accuracy: 0.8274 - val_loss: 0.5128 - learning_rate: 0.0010
Epoch 18/50


352/352  **0s** 16ms/step - accuracy: 0.8567 - loss: 0.4174
Epoch 18: val_accuracy did not improve from 0.82940


352/352  **6s** 17ms/step - accuracy: 0.8567 - loss: 0.4174 - val_accuracy: 0.8138 - val_loss: 0.5673 - learning_rate: 0.0010
Epoch 19/50


352/352  **0s** 17ms/step - accuracy: 0.8595 - loss: 0.4067


Epoch 19: val_accuracy improved from 0.82940 to 0.83540, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **6s** 17ms/step - accuracy: 0.8596 - loss: 0.4067 - val_accuracy: 0.8354 - val_loss: 0.5184 - learning_rate: 0.0010
Epoch 20/50

349/352  **0s** 17ms/step - accuracy: 0.8679 - loss: 0.3866
Epoch 20: val_accuracy did not improve from 0.83540

352/352  **6s** 18ms/step - accuracy: 0.8679 - loss: 0.3865 - val_accuracy: 0.8304 - val_loss: 0.6060 - learning_rate: 0.0010
Epoch 21/50


351/352  **0s** 17ms/step - accuracy: 0.8719 - loss: 0.3756
Epoch 21: val_accuracy did not improve from 0.83540


352/352  **6s** 17ms/step - accuracy: 0.8719 - loss: 0.3756 - val_accuracy: 0.8144 - val_loss: 0.6033 - learning_rate: 0.0010
Epoch 22/50

349/352  **0s** 17ms/step - accuracy: 0.8719 - loss: 0.3677

Epoch 22: val_accuracy improved from 0.83540 to 0.83980, saving model to best_custom_cnn_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **6s** 18ms/step - accuracy: 0.8720 - loss: 0.3676 - val_accuracy: 0.8398 - val_loss: 0.5183 - learning_rate: 0.0010
Epoch 23/50

351/352  **0s** 17ms/step - accuracy: 0.8888 - loss: 0.3171


Epoch 23: val_accuracy improved from 0.83980 to 0.85800, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  **6s** 18ms/step - accuracy: 0.8888 - loss: 0.3170 - val_accuracy: 0.8580 - val_loss: 0.4724 - learning_rate: 2.0000e-04
Epoch 24/50

351/352  **0s** 17ms/step - accuracy: 0.9029 - loss: 0.2776


Epoch 24: val_accuracy did not improve from 0.85800


352/352  **6s** 18ms/step - accuracy: 0.9030 - loss: 0.2775 - val_accuracy: 0.8580 - val_loss: 0.4704 - learning_rate: 2.0000e-04
Epoch 25/50

351/352  **0s** 17ms/step - accuracy: 0.9084 - loss: 0.2674

Epoch 25: val_accuracy improved from 0.85800 to 0.86100, saving model to best_custom_cnn_model.h5







WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  **6s** 18ms/step - accuracy: 0.9084 - loss: 0.2673 - val_accuracy: 0.8610 - val_loss: 0.4680 - learning_rate: 2.0000e-04
Epoch 26/50


















349/352  **0s** 17ms/step - accuracy: 0.9097 - loss: 0.2549





Epoch 26: val_accuracy improved from 0.86100 to 0.86160, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  **6s** 18ms/step - accuracy: 0.9097 - loss: 0.2547 - val_accuracy: 0.8616 - val_loss: 0.4728 - learning_rate: 2.0000e-04
Epoch 27/50
350/352  **0s** 17ms/step - accuracy: 0.9161 - loss: 0.2408
Epoch 27: val_accuracy did not improve from 0.86160
352/352  **6s** 18ms/step - accuracy: 0.9162 - loss: 0.2408 - val_accuracy: 0.8602 - val_loss: 0.4760 - learning_rate: 2.0000e-04
Epoch 28/50
350/352  **0s** 17ms/step - accuracy: 0.9177 - loss: 0.2338
Epoch 28: val_accuracy did not improve from 0.86160
352/352  **6s** 18ms/step - accuracy: 0.9177 - loss: 0.2337 - val_accuracy: 0.8578 - val_loss: 0.4910 - learning_rate: 2.0000e-04
Epoch 29/50
350/352  **0s** 17ms/step - accuracy: 0.9214 - loss: 0.2205
Epoch 29: val_accuracy improved from 0.86160 to 0.86440, saving model to best_custom_cnn_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  **6s** 18ms/step - accuracy: 0.9214 - loss: 0.2204 - val_accuracy: 0.8644 - val_loss: 0.4852 - learning_rate: 2.0000e-04
Epoch 30/50
351/352  **0s** 17ms/step - accuracy: 0.9237 - loss: 0.2209
Epoch 30: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9237 - loss: 0.2209 - val_accuracy: 0.8616 - val_loss: 0.4828 - learning_rate: 2.0000e-04
Epoch 31/50
351/352  **0s** 16ms/step - accuracy: 0.9253 - loss: 0.2126
Epoch 31: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9253 - loss: 0.2126 - val_accuracy: 0.8642 - val_loss: 0.4916 - learning_rate: 1.0000e-04
Epoch 32/50
351/352  **0s** 16ms/step - accuracy: 0.9270 - loss: 0.2080
Epoch 32: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9270 - loss: 0.2079 - val_accuracy: 0.8634 - val_loss: 0.4851 - learning_rate: 1.0000e-04
Epoch 33/50
352/352  **0s** 16ms/step - accuracy: 0.9291 - loss: 0.1962
Epoch 33: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9291 - loss: 0.1961 - val_accuracy: 0.8638 - val_loss: 0.4912 - learning_rate: 1.0000e-04
Epoch 34/50
350/352  **0s** 16ms/step - accuracy: 0.9315 - loss: 0.1948
Epoch 34: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9315 - loss: 0.1947 - val_accuracy: 0.8602 - val_loss: 0.4953 - learning_rate: 1.0000e-04
Epoch 35/50
351/352  **0s** 16ms/step - accuracy: 0.9333 - loss: 0.1911
Epoch 35: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9333 - loss: 0.1910 - val_accuracy: 0.8606 - val_loss: 0.4897 - learning_rate: 1.0000e-04
Epoch 36/50
350/352  **0s** 16ms/step - accuracy: 0.9339 - loss: 0.1846
Epoch 36: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9339 - loss: 0.1845 - val_accuracy: 0.8604 - val_loss: 0.4911 - learning_rate: 1.0000e-04
Epoch 37/50
352/352  **0s** 16ms/step - accuracy: 0.9362 - loss: 0.1829
Epoch 37: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9362 - loss: 0.1829 - val_accuracy: 0.8616 - val_loss: 0.491

9 - learning_rate: 1.0000e-04
 Epoch 38/50
349/352  **0s** 16ms/step - accuracy: 0.9374 - loss: 0.1767
 Epoch 38: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9374 - loss: 0.1767 - val_accuracy: 0.8630 - val_loss: 0.501
 7 - learning_rate: 1.0000e-04
 Epoch 39/50
350/352  **0s** 16ms/step - accuracy: 0.9375 - loss: 0.1803
 Epoch 39: val_accuracy did not improve from 0.86440
352/352  **6s** 17ms/step - accuracy: 0.9375 - loss: 0.1802 - val_accuracy: 0.8604 - val_loss: 0.505
 1 - learning_rate: 1.0000e-04
 Custom CNN training completed in 257.18 seconds

Training AlexNet

```
In [9]: print("Training AlexNet...")
start_time = time.time()

# Train AlexNet
alexnet_history = alexnet_model.fit(
    x_train_normalized, y_train_categorical,
    batch_size=BATCH_SIZE,
    epochs=EPOCHS,
    validation_split=VALIDATION_SPLIT,
    callbacks=create_callbacks('alexnet'),
    verbose=1
)


alexnet_training_time = time.time() - start_time
print(f"AlexNet training completed in {alexnet_training_time:.2f} seconds")
```

Training AlexNet...

Epoch 1/50

351/352  **0s** 74ms/step - accuracy: 0.2323 - loss: 2.4036

2025-07-28 20:27:04.625738: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_2024', 60 bytes spill stores, 60 bytes spill loads

352/352  **0s** 107ms/step - accuracy: 0.2325 - loss: 2.4020

2025-07-28 20:27:15.896641: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 532 bytes spill stores, 532 bytes spill loads

2025-07-28 20:27:16.015756: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 3148 bytes spill stores, 3124 bytes spill loads

2025-07-28 20:27:16.055718: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 488 bytes spill stores, 488 bytes spill loads

2025-07-28 20:27:16.146785: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 4 bytes spill stores, 4 bytes spill loads

2025-07-28 20:27:16.224192: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 820 bytes spill stores, 820 bytes spill loads

2025-07-28 20:27:16.289605: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160_0', 136 bytes spill stores, 136 bytes spill loads

2025-07-28 20:27:16.674302: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160_0', 1604 bytes spill stores, 1552 bytes spill loads

2025-07-28 20:27:16.722740: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 8 bytes spill stores, 8 bytes spill loads

2025-07-28 20:27:16.815072: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 4 bytes spill stores, 4 bytes spill loads

2025-07-28 20:27:16.863591: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 128 bytes spill stores, 128 bytes spill loads

2025-07-28 20:27:16.896665: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 64 bytes spill stores, 64 bytes spill loads

2025-07-28 20:27:17.022873: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 76 bytes spill stores, 76 bytes spill loads

2025-07-28 20:27:17.181230: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 4 bytes spill stores, 4 bytes spill loads

2025-07-28 20:27:17.298295: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 896 bytes spill stores, 896 bytes spill loads

2025-07-28 20:27:17.470751: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 4284 bytes spill stores, 4260 bytes spill loads

2025-07-28 20:27:19.767977: E external/local_xla/xla/service/slow_operation_alarm.cc:73] Trying algorithm eng0{} for conv %cudnn-conv-bias-activation.16 = (f32[128,256,16,16]{3,2,1,0}, u8[0]{0}) custom-call(f32[128,96,16,16]{3,2,1,0} %bitcast.729, f32[256,96,5,5]{3,2,1,0} %bitcast.736, f32[256]{0} %bitcast.738), window={size=5x5 pad=2_2x2_2}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn\$convBiasActivationForward", metadata={op_type="Conv2D" op_name="sequential_1_1/conv2d_7_1/convolution" source_file="/home/abhijit/miniconda3/envs/tf-env/lib/python3.12/site-packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_config":{"conv_result_scale":1,"activation_mode":"kRelu","side_input_scale":0,"leakyrelu_alpha":0},"force_earliest_schedule":false} is taking a while...

2025-07-28 20:27:20.364799: E external/local_xla/xla/service/slow_operation_alarm.cc:140] The operation took 1.59693629s

Trying algorithm eng0{} for conv %cudnn-conv-bias-activation.16 = (f32[128,256,16,16]{3,2,1,0}, u8[0]{0}) custom-call(f32[128,96,16,16]{3,2,1,0} %bitcast.729, f32[256,96,5,5]{3,2,1,0} %bitcast.736, f32[256]{0} %bitcast.738), window={size=5x5 pad=2_2x2_2}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn\$convBiasActivationForward", metadata={op_type="Conv2D" op_name="sequential_1_1/conv2d_7_1/convolution" source_file="/home/abhijit/miniconda3/envs/tf-env/lib/python3.12/site-packages/tensorflow/python/framework/ops.py" source_line=1200}, backend_config={"operation_queue_id":"0","wait_on_operation_queues":[],"cudnn_conv_backend_config":{"conv_result_scale":1,"activation_mode":"kRelu","side_input_scale":0,"leakyrelu_alpha":0},"force_earliest_schedule":false} is taking a while...

2025-07-28 20:27:24.339164: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 56 bytes spill stores, 56 bytes spill loads

2025-07-28 20:27:24.376780: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 408 bytes spill stores, 408 bytes spill loads

2025-07-28 20:27:24.590444: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 4 bytes spill stores, 4 bytes spill loads


2025-07-28 20:27:24.836074: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 844 bytes spill stores, 844 bytes spill loads

2025-07-28 20:27:24.996506: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 100 bytes spill stores, 100 bytes spill loads

2025-07-28 20:27:25.056185: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160_0', 3392 bytes spill stores, 3120 bytes spill loads

Epoch 1: val_accuracy improved from -inf to 0.22380, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **64s** 140ms/step - accuracy: 0.2328 - loss: 2.4005 - val_accuracy: 0.2238 - val_loss: 2.4434 - learning_rate: 0.0010

Epoch 2/50

351/352  **0s** 79ms/step - accuracy: 0.4797 - loss: 1.4224

Epoch 2: val_accuracy improved from 0.22380 to 0.41820, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **29s** 82ms/step - accuracy: 0.4799 - loss: 1.4220 - val_accuracy: 0.4182 - val_loss: 1.6395 - learning_rate: 0.0010

Epoch 3/50


352/352  **0s** 81ms/step - accuracy: 0.5732 - loss: 1.1931

Epoch 3: val_accuracy improved from 0.41820 to 0.48860, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **30s** 84ms/step - accuracy: 0.5732 - loss: 1.1930 - val_accuracy: 0.4886 - val_loss: 1.4592 - learning_rate: 0.0010

Epoch 4/50


352/352  **0s** 81ms/step - accuracy: 0.6336 - loss: 1.0468

Epoch 4: val_accuracy improved from 0.48860 to 0.53160, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **30s** 84ms/step - accuracy: 0.6336 - loss: 1.0467 - val_accuracy: 0.5316 - val_loss: 1.4276 - learning_rate: 0.0010

Epoch 5/50


352/352  **0s** 78ms/step - accuracy: 0.6816 - loss: 0.9027

Epoch 5: val_accuracy improved from 0.53160 to 0.54560, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **29s** 82ms/step - accuracy: 0.6816 - loss: 0.9026 - val_accuracy: 0.5456 - val_loss: 1.2956 - learning_rate: 0.0010

Epoch 6/50


351/352  **0s** 79ms/step - accuracy: 0.7200 - loss: 0.8033

Epoch 6: val_accuracy improved from 0.54560 to 0.61180, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **29s** 82ms/step - accuracy: 0.7201 - loss: 0.8032 - val_accuracy: 0.6118 - val_loss: 1.1090 - learning_rate: 0.0010


Epoch 7/50


351/352  **0s** 77ms/step - accuracy: 0.7431 - loss: 0.7411


Epoch 7: val_accuracy improved from 0.61180 to 0.67320, saving model to best_alexnet_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  28s 81ms/step - accuracy: 0.7432 - loss: 0.7410 - val_accuracy: 0.6732 - val_loss: 0.9536 - learning_rate: 0.0010
Epoch 8/50


352/352  0s 78ms/step - accuracy: 0.7690 - loss: 0.6664
Epoch 8: val_accuracy did not improve from 0.67320


352/352  28s 80ms/step - accuracy: 0.7690 - loss: 0.6664 - val_accuracy: 0.4462 - val_loss: 1.8438 - learning_rate: 0.0010
Epoch 9/50


352/352  0s 77ms/step - accuracy: 0.7878 - loss: 0.6154
Epoch 9: val_accuracy did not improve from 0.67320


352/352  28s 79ms/step - accuracy: 0.7878 - loss: 0.6154 - val_accuracy: 0.5324 - val_loss: 1.4436 - learning_rate: 0.0010
Epoch 10/50


351/352  0s 81ms/step - accuracy: 0.8059 - loss: 0.5608
Epoch 10: val_accuracy did not improve from 0.67320


352/352  29s 83ms/step - accuracy: 0.8059 - loss: 0.5607 - val_accuracy: 0.6706 - val_loss: 1.0698 - learning_rate: 0.0010
Epoch 11/50

351/352  0s 79ms/step - accuracy: 0.8226 - loss: 0.5160
Epoch 11: val_accuracy did not improve from 0.67320

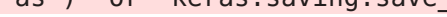
352/352  28s 81ms/step - accuracy: 0.8226 - loss: 0.5159 - val_accuracy: 0.6592 - val_loss: 1.1102 - learning_rate: 0.0010
Epoch 12/50


352/352  0s 79ms/step - accuracy: 0.8341 - loss: 0.4824
Epoch 12: val_accuracy did not improve from 0.67320

352/352  28s 81ms/step - accuracy: 0.8341 - loss: 0.4823 - val_accuracy: 0.6520 - val_loss: 1.0895 - learning_rate: 0.0010
Epoch 13/50


352/352  0s 80ms/step - accuracy: 0.8838 - loss: 0.3427
Epoch 13: val_accuracy improved from 0.67320 to 0.75220, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  29s 83ms/step - accuracy: 0.8838 - loss: 0.3424 - val_accuracy: 0.7522 - val_loss: 0.9102 - learning_rate: 2.0000e-04
Epoch 14/50

351/352  0s 78ms/step - accuracy: 0.9426 - loss: 0.1688
Epoch 14: val_accuracy improved from 0.75220 to 0.75820, saving model to best_alexnet_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

352/352  29s 81ms/step - accuracy: 0.9426 - loss: 0.1686 - val_accuracy: 0.7582 - val_loss: 1.1199 - learning_rate: 2.0000e-04


Epoch 15/50

351/352  0s 77ms/step - accuracy: 0.9724 - loss: 0.0824


Epoch 15: val_accuracy did not improve from 0.75820

352/352  28s 79ms/step - accuracy: 0.9724 - loss: 0.0823 - val_accuracy: 0.7496 - val_loss: 1.4581 - learning_rate: 2.0000e-04

Epoch 16/50

351/352  0s 79ms/step - accuracy: 0.9835 - loss: 0.0511


Epoch 16: val_accuracy did not improve from 0.75820

352/352  28s 81ms/step - accuracy: 0.9835 - loss: 0.0511 - val_accuracy: 0.7556 - val_loss: 1.6869 - learning_rate: 2.0000e-04


Epoch 17/50

351/352  0s 77ms/step - accuracy: 0.9877 - loss: 0.0378


Epoch 17: val_accuracy did not improve from 0.75820

352/352  28s 78ms/step - accuracy: 0.9877 - loss: 0.0377 - val_accuracy: 0.7536 - val_loss: 1.7954 - learning_rate: 2.0000e-04


Epoch 18/50

351/352  0s 77ms/step - accuracy: 0.9885 - loss: 0.0344

Epoch 18: val_accuracy did not improve from 0.75820


352/352  28s 79ms/step - accuracy: 0.9885 - loss: 0.0344 - val_accuracy: 0.7582 - val_loss: 1.7593 - learning_rate: 2.0000e-04

Epoch 19/50

351/352  0s 77ms/step - accuracy: 0.9928 - loss: 0.0225

Epoch 19: val_accuracy improved from 0.75820 to 0.76780, saving model to best_alexnet_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  28s 80ms/step - accuracy: 0.9928 - loss: 0.0225 - val_accuracy: 0.7678 - val_loss: 1.8207 - learning_rate: 1.0000e-04

Epoch 20/50

351/352  0s 78ms/step - accuracy: 0.9958 - loss: 0.0133

Epoch 20: val_accuracy improved from 0.76780 to 0.76920, saving model to best_alexnet_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **29s** 81ms/step - accuracy: 0.9959 - loss: 0.0133 - val_accuracy: 0.7692 - val_loss: 1.9272 - learning_rate: 1.0000e-04
Epoch 21/50


351/352  **0s** 78ms/step - accuracy: 0.9985 - loss: 0.0060


Epoch 21: val_accuracy improved from 0.76920 to 0.77240, saving model to best_alexnet_model.h5


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.


352/352  **29s** 81ms/step - accuracy: 0.9985 - loss: 0.0060 - val_accuracy: 0.7724 - val_loss: 1.9747 - learning_rate: 1.0000e-04
Epoch 22/50


351/352  **0s** 79ms/step - accuracy: 0.9983 - loss: 0.0060
Epoch 22: val_accuracy did not improve from 0.77240


352/352  **29s** 81ms/step - accuracy: 0.9983 - loss: 0.0060 - val_accuracy: 0.7706 - val_loss: 2.0802 - learning_rate: 1.0000e-04
Epoch 23/50


352/352  **0s** 81ms/step - accuracy: 0.9985 - loss: 0.0046
Epoch 23: val_accuracy did not improve from 0.77240


352/352  **29s** 83ms/step - accuracy: 0.9985 - loss: 0.0046 - val_accuracy: 0.7698 - val_loss: 2.1056 - learning_rate: 1.0000e-04
Epoch 24/50


351/352  **0s** 85ms/step - accuracy: 0.9989 - loss: 0.0041
Epoch 24: val_accuracy did not improve from 0.77240


352/352  **31s** 87ms/step - accuracy: 0.9989 - loss: 0.0041 - val_accuracy: 0.7658 - val_loss: 2.3512 - learning_rate: 1.0000e-04
Epoch 25/50


352/352  **0s** 80ms/step - accuracy: 0.9988 - loss: 0.0037
Epoch 25: val_accuracy did not improve from 0.77240


352/352  **29s** 81ms/step - accuracy: 0.9988 - loss: 0.0037 - val_accuracy: 0.7712 - val_loss: 2.4086 - learning_rate: 1.0000e-04
Epoch 26/50


352/352  **0s** 79ms/step - accuracy: 0.9975 - loss: 0.0086
Epoch 26: val_accuracy did not improve from 0.77240


352/352  **29s** 81ms/step - accuracy: 0.9975 - loss: 0.0086 - val_accuracy: 0.7578 - val_loss: 2.4284 - learning_rate: 1.0000e-04
Epoch 27/50


351/352  **0s** 74ms/step - accuracy: 0.9971 - loss: 0.0105
Epoch 27: val_accuracy did not improve from 0.77240





352/352  **27s** 76ms/step - accuracy: 0.9971 - loss: 0.0105 - val_accuracy: 0.7696 - val_loss: 2.3189 - learning_rate: 1.0000e-04
Epoch 28/50

351/352  **0s** 64ms/step - accuracy: 0.9989 - loss: 0.0040
Epoch 28: val_accuracy did not improve from 0.77240

352/352  **23s** 65ms/step - accuracy: 0.9989 - loss: 0.0040 - val_accuracy: 0.7712 - val_loss: 2.3914 - learning_rate: 1.0000e-04
Epoch 29/50

351/352  **0s** 57ms/step - accuracy: 0.9979 - loss: 0.0067
Epoch 29: val_accuracy did not improve from 0.77240

352/352  **20s** 58ms/step - accuracy: 0.9979 - loss: 0.0066 - val_accuracy: 0.7720 - val_loss: 2.38

63 - learning_rate: 1.0000e-04
 Epoch 30/50
 351/352  0s 56ms/step - accuracy: 0.9985 - loss: 0.0059
 Epoch 30: val_accuracy did not improve from 0.77240
 352/352  20s 58ms/step - accuracy: 0.9985 - loss: 0.0059 - val_accuracy: 0.7698 - val_loss: 2.61
 28 - learning_rate: 1.0000e-04
 Epoch 31/50
 351/352  0s 57ms/step - accuracy: 0.9989 - loss: 0.0040
 Epoch 31: val_accuracy did not improve from 0.77240
 352/352  21s 58ms/step - accuracy: 0.9989 - loss: 0.0040 - val_accuracy: 0.7684 - val_loss: 2.69
 89 - learning_rate: 1.0000e-04
 AlexNet training completed in 892.38 seconds

Model Evaluation and Testing

```
In [10]: # Evaluate both models on test set
print("Evaluating Custom CNN on test set...")
custom_cnn_test_loss, custom_cnn_test_accuracy = custom_cnn.evaluate(
    x_test_normalized, y_test_categorical, verbose=0
)

print("Evaluating AlexNet on test set...")
alexnet_test_loss, alexnet_test_accuracy = alexnet_model.evaluate(
    x_test_normalized, y_test_categorical, verbose=0
)

# Generate predictions for detailed analysis
custom_cnn_predictions = custom_cnn.predict(x_test_normalized)
alexnet_predictions = alexnet_model.predict(x_test_normalized)

custom_cnn_pred_classes = np.argmax(custom_cnn_predictions, axis=1)
alexnet_pred_classes = np.argmax(alexnet_predictions, axis=1)
true_classes = np.argmax(y_test_categorical, axis=1)

print(f"\nTest Results:")
print(f"Custom CNN - Test Accuracy: {custom_cnn_test_accuracy:.4f} ({custom_cnn_test_accuracy*100:.2f}%)")
print(f"AlexNet - Test Accuracy: {alexnet_test_accuracy:.4f} ({alexnet_test_accuracy*100:.2f}%)")
```

Evaluating Custom CNN on test set...

```
2025-07-28 20:41:14.684317: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_207', 4 bytes spill stores, 4 bytes spill loads
```

Evaluating AlexNet on test set...

2025-07-28 20:41:17.694746: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 64 bytes spill stores, 64 bytes spill loads

2025-07-28 20:41:17.853248: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 488 bytes spill stores, 488 bytes spill loads

2025-07-28 20:41:17.977578: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 532 bytes spill stores, 532 bytes spill loads

2025-07-28 20:41:17.996921: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160_0', 48 bytes spill stores, 48 bytes spill loads

2025-07-28 20:41:17.999872: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 820 bytes spill stores, 820 bytes spill loads

2025-07-28 20:41:18.276553: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 4 bytes spill stores, 4 bytes spill loads

2025-07-28 20:41:18.323096: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 3148 bytes spill stores, 3124 bytes spill loads

2025-07-28 20:41:18.546823: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 76 bytes spill stores, 76 bytes spill loads

2025-07-28 20:41:18.614435: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 4284 bytes spill stores, 4260 bytes spill loads

2025-07-28 20:41:18.722345: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 896 bytes spill stores, 896 bytes spill loads

```
2025-07-28 20:41:21.707077: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 64 bytes spill stores, 64 bytes spill loads
```

```
2025-07-28 20:41:22.144993: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 64 bytes spill stores, 64 bytes spill loads
```

```
2025-07-28 20:41:22.170171: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 532 bytes spill stores, 532 bytes spill loads
```

```
2025-07-28 20:41:22.380889: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_160', 3148 bytes spill stores, 3124 bytes spill loads
```

```
2025-07-28 20:41:22.429729: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_167', 896 bytes spill stores, 896 bytes spill loads
```

```
313/313 ————— 1s 2ms/step
313/313 ————— 2s 5ms/step
```

Test Results:

Custom CNN - Test Accuracy: 0.8574 (85.74%)

AlexNet - Test Accuracy: 0.7533 (75.33%)

Training History Visualization

```
In [11]: # Plot training history
def plot_training_history(custom_history, alexnet_history):
    fig, axes = plt.subplots(2, 2, figsize=(15, 10))

    # Accuracy plots
    axes[0, 0].plot(custom_history.history['accuracy'], label='Custom CNN Train', color='blue')
    axes[0, 0].plot(custom_history.history['val_accuracy'], label='Custom CNN Val', color='lightblue')
    axes[0, 0].set_title('Custom CNN - Accuracy')
    axes[0, 0].set_xlabel('Epoch')
    axes[0, 0].set_ylabel('Accuracy')
```



```

axes[0, 0].legend()
axes[0, 0].grid(True)

axes[0, 1].plot(alexnet_history.history['accuracy'], label='AlexNet Train', color='red')
axes[0, 1].plot(alexnet_history.history['val_accuracy'], label='AlexNet Val', color='lightcoral')
axes[0, 1].set_title('AlexNet - Accuracy')
axes[0, 1].set_xlabel('Epoch')
axes[0, 1].set_ylabel('Accuracy')
axes[0, 1].legend()
axes[0, 1].grid(True)

# Loss plots
axes[1, 0].plot(custom_history.history['loss'], label='Custom CNN Train', color='blue')
axes[1, 0].plot(custom_history.history['val_loss'], label='Custom CNN Val', color='lightblue')
axes[1, 0].set_title('Custom CNN - Loss')
axes[1, 0].set_xlabel('Epoch')
axes[1, 0].set_ylabel('Loss')
axes[1, 0].legend()
axes[1, 0].grid(True)

axes[1, 1].plot(alexnet_history.history['loss'], label='AlexNet Train', color='red')
axes[1, 1].plot(alexnet_history.history['val_loss'], label='AlexNet Val', color='lightcoral')
axes[1, 1].set_title('AlexNet - Loss')
axes[1, 1].set_xlabel('Epoch')
axes[1, 1].set_ylabel('Loss')
axes[1, 1].legend()
axes[1, 1].grid(True)





plt.tight_layout()
plt.show()

```

```

plot_training_history(custom_cnn_history, alexnet_history)

```



Comparative Analysis

```

In [12]: # Comparative accuracy plot
plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
epochs_range = range(1, len(custom_cnn_history.history['accuracy']) + 1)
plt.plot(epochs_range, custom_cnn_history.history['accuracy'], 'b-', label='Custom CNN Train')
plt.plot(epochs_range, custom_cnn_history.history['val_accuracy'], 'b--', label='Custom CNN Val')

epochs_range_alex = range(1, len(alexnet_history.history['accuracy']) + 1)
plt.plot(epochs_range_alex, alexnet_history.history['accuracy'], 'r-', label='AlexNet Train')
plt.plot(epochs_range_alex, alexnet_history.history['val_accuracy'], 'r--', label='AlexNet Val')

plt.title('Training and Validation Accuracy Comparison')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)



plt.subplot(1, 2, 2)
models = ['Custom CNN', 'AlexNet']
test_accuracies = [custom_cnn_test_accuracy, alexnet_test_accuracy]
colors = ['blue', 'red']

bars = plt.bar(models, test_accuracies, color=colors, alpha=0.7)
plt.title('Test Accuracy Comparison')
plt.ylabel('Test Accuracy')
plt.ylim(0, 1)

# Add value labels on bars
for bar, acc in zip(bars, test_accuracies):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.01,
             f'{acc:.3f}', ha='center', va='bottom')

plt.tight_layout()
plt.show()

```



Confusion Matrix and Classification Reports

```
In [13]: # Generate classification reports
print("Classification Report - Custom CNN:")
print(classification_report(true_classes, custom_cnn_pred_classes, target_names=class_names))

print("\nClassification Report - AlexNet:")
print(classification_report(true_classes, alexnet_pred_classes, target_names=class_names))
```

Classification Report - Custom CNN:

	precision	recall	f1-score	support
airplane	0.88	0.86	0.87	1000
automobile	0.93	0.94	0.93	1000
bird	0.80	0.79	0.80	1000
cat	0.74	0.70	0.72	1000
deer	0.83	0.86	0.85	1000
dog	0.79	0.79	0.79	1000
frog	0.89	0.89	0.89	1000
horse	0.89	0.89	0.89	1000
ship	0.90	0.93	0.91	1000
truck	0.91	0.91	0.91	1000
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

Classification Report - AlexNet:



	precision	recall	f1-score	support
airplane	0.76	0.80	0.78	1000
automobile	0.85	0.89	0.87	1000
bird	0.64	0.65	0.65	1000
cat	0.55	0.57	0.56	1000
deer	0.71	0.70	0.71	1000
dog	0.66	0.64	0.65	1000
frog	0.83	0.81	0.82	1000
horse	0.82	0.79	0.80	1000
ship	0.87	0.86	0.86	1000
truck	0.85	0.81	0.83	1000
accuracy			0.75	10000
macro avg	0.75	0.75	0.75	10000
weighted avg	0.75	0.75	0.75	10000

```
In [14]: # Plot confusion matrices
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
```

```
# Custom CNN confusion matrix
cm_custom = confusion_matrix(true_classes, custom_cnn_pred_classes)
sns.heatmap(cm_custom, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names, ax=axes[0])
axes[0].set_title(f'Custom CNN Confusion Matrix\nTest Accuracy: {custom_cnn_test_accuracy:.3f}')
axes[0].set_xlabel('Predicted')
axes[0].set_ylabel('Actual')

# AlexNet confusion matrix
cm_alexnet = confusion_matrix(true_classes, alexnet_pred_classes)
sns.heatmap(cm_alexnet, annot=True, fmt='d', cmap='Reds',
            xticklabels=class_names, yticklabels=class_names, ax=axes[1])
axes[1].set_title(f'AlexNet Confusion Matrix\nTest Accuracy: {alexnet_test_accuracy:.3f}')
axes[1].set_xlabel('Predicted')
axes[1].set_ylabel('Actual')

plt.tight_layout()
plt.show()
```



Model Complexity Analysis

```
In [15]: # Calculate model parameters and complexity
def count_parameters(model):
    return model.count_params()

custom_cnn_params = count_parameters(custom_cnn)
alexnet_params = count_parameters(alexnet_model)

print("Model Complexity Analysis:")
print(f"Custom CNN:")
print(f" - Total Parameters: {custom_cnn_params:,}")
print(f" - Training Time: {custom_cnn_training_time:.2f} seconds")
print(f" - Test Accuracy: {custom_cnn_test_accuracy:.4f}")

print(f"\nAlexNet:")
print(f" - Total Parameters: {alexnet_params:,}")
print(f" - Training Time: {alexnet_training_time:.2f} seconds")
print(f" - Test Accuracy: {alexnet_test_accuracy:.4f}")

print(f"\nParameter Ratio: AlexNet has {alexnet_params/custom_cnn_params:.1f}x more parameters than Custom CNN")
```

Model Complexity Analysis:

Custom CNN:

- Total Parameters: 1,472,938
- Training Time: 257.18 seconds
- Test Accuracy: 0.8574

AlexNet:

- Total Parameters: 37,324,554
- Training Time: 892.38 seconds
- Test Accuracy: 0.7533

Parameter Ratio: AlexNet has 25.3x more parameters than Custom CNN

Sample Predictions Visualization

```
In [16]: # Visualize some predictions
```

```

def plot_predictions(images, true_labels, custom_predictions, alexnet_predictions, class_names, num_samples=12):
    plt.figure(figsize=(15, 10))

    for i in range(num_samples):
        plt.subplot(3, 4, i + 1)
        plt.imshow(images[i])

        true_class = class_names[true_labels[i]]
        custom_pred_class = class_names[custom_predictions[i]]
        alexnet_pred_class = class_names[alexnet_predictions[i]]

        # Color code: green if both correct, red if both wrong, yellow if mixed
        custom_correct = custom_predictions[i] == true_labels[i]
        alexnet_correct = alexnet_predictions[i] == true_labels[i]

        if custom_correct and alexnet_correct:
            color = 'green'
        elif not custom_correct and not alexnet_correct:
            color = 'red'
        else:
            color = 'orange'

        plt.title(f'True: {true_class}\nCustom: {custom_pred_class}\nAlexNet: {alexnet_pred_class}',
                  fontsize=8, color=color)
        plt.axis('off')

    plt.suptitle('Sample Predictions Comparison\n(Green: Both Correct, Red: Both Wrong, Orange: Mixed)',
                  fontsize=12)
    plt.tight_layout()
    plt.show()

# Show some random predictions
random_indices = np.random.choice(len(x_test), 12, replace=False)
plot_predictions(x_test[random_indices], true_classes[random_indices],
                 custom_cnn_pred_classes[random_indices], alexnet_pred_classes[random_indices],
                 class_names)

```


Sample Predictions Comparison (Green: Both Correct, Red: Both Wrong, Orange: Mixed)


True: bird
Custom: bird
AlexNet: bird




True: automobile
Custom: automobile
AlexNet: automobile



True: dog
Custom: cat
AlexNet: cat



True: ship
Custom: ship
AlexNet: ship



True: truck
Custom: cat
AlexNet: truck



True: cat
Custom: cat
AlexNet: cat



True: ship
Custom: ship
AlexNet: ship



True: truck
Custom: truck
AlexNet: truck



True: airplane
Custom: airplane
AlexNet: airplane



True: automobile
Custom: truck
AlexNet: automobile



True: dog
Custom: dog
AlexNet: dog



True: automobile
Custom: automobile
AlexNet: automobile



Final Performance Summary and Analysis

```
In [17]: print("="*80)
print("CIFAR-10 CLASSIFICATION - FINAL PERFORMANCE SUMMARY")
print("="*80)

print(f"\n📊 TEST ACCURACY RESULTS:")
print(f"Custom CNN:  {custom_cnn_test_accuracy:.4f} ({custom_cnn_test_accuracy*100:.2f}%)")
print(f"AlexNet:      {alexnet_test_accuracy:.4f} ({alexnet_test_accuracy*100:.2f}%)")

if custom_cnn_test_accuracy > alexnet_test_accuracy:
    winner = "Custom CNN"
    difference = custom_cnn_test_accuracy - alexnet_test_accuracy
else:
    winner = "AlexNet"
    difference = alexnet_test_accuracy - custom_cnn_test_accuracy

print(f"\n🏆 WINNER: {winner} (by {difference*100:.2f} percentage points)")

print(f"\n⚙️ MODEL COMPLEXITY:")
print(f"Custom CNN Parameters:  {custom_cnn_params:,}")
print(f"AlexNet Parameters:     {alexnet_params:,}")
print(f"Parameter Ratio:        {alexnet_params/custom_cnn_params:.1f}x")

print(f"\n⌚ TRAINING TIME:")
print(f"Custom CNN:  {custom_cnn_training_time:.2f} seconds")
print(f"AlexNet:     {alexnet_training_time:.2f} seconds")

print(f"\n📝 OBSERVATIONS AND ANALYSIS:")
print(f"\n1. ACCURACY COMPARISON:")
if custom_cnn_test_accuracy > alexnet_test_accuracy:
    print(f"    • Custom CNN outperformed AlexNet by {difference*100:.2f}%")
    print(f"    • This suggests that the custom architecture is better suited for CIFAR-10")
else:
    print(f"    • AlexNet outperformed Custom CNN by {difference*100:.2f}%")
    print(f"    • This demonstrates the power of the proven AlexNet architecture")

print(f"\n2. MODEL EFFICIENCY:")
custom_efficiency = custom_cnn_test_accuracy / (custom_cnn_params / 1000000)
```

```

alexnet_efficiency = alexnet_test_accuracy / (alexnet_params / 1000000)
print(f"    • Custom CNN Efficiency: {custom_efficiency:.3f} (accuracy per million parameters)")
print(f"    • AlexNet Efficiency: {alexnet_efficiency:.3f} (accuracy per million parameters)")

if custom_efficiency > alexnet_efficiency:
    print(f"    • Custom CNN is more parameter-efficient")
else:
    print(f"    • AlexNet achieves better accuracy despite having more parameters")

print(f"\n3. ARCHITECTURE INSIGHTS:")
print(f"    • Custom CNN uses modern techniques: Batch Normalization, Dropout")
print(f"    • AlexNet adapted from 224x224 to 32x32 input size")
print(f"    • Both models benefit from data normalization and proper regularization")

print(f"\n4. PRACTICAL CONSIDERATIONS:")
if custom_cnn_training_time < alexnet_training_time:
    print(f"    • Custom CNN trains {alexnet_training_time/custom_cnn_training_time:.1f}x faster")
else:
    print(f"    • AlexNet trains {custom_cnn_training_time/alexnet_training_time:.1f}x faster")
print(f"    • Model size affects deployment and inference speed")
print(f"    • Custom architectures can be optimized for specific datasets")

print(f"\n" + "="*80)
print(f"CONCLUSION: The {winner} model provides the best performance for CIFAR-10 classification")
print(f"in terms of test accuracy, achieving {max(custom_cnn_test_accuracy, alexnet_test_accuracy)*100:.2f}% accuracy")
print(f"="*80)

```

CIFAR-10 CLASSIFICATION - FINAL PERFORMANCE SUMMARY

TEST ACCURACY RESULTS:

Custom CNN: 0.8574 (85.74%)

AlexNet: 0.7533 (75.33%)

 WINNER: Custom CNN (by 10.41 percentage points)

MODEL COMPLEXITY:

Custom CNN Parameters: 1,472,938

AlexNet Parameters: 37,324,554

Parameter Ratio: 25.3x

TRAINING TIME:

Custom CNN: 257.18 seconds

AlexNet: 892.38 seconds

OBSERVATIONS AND ANALYSIS:

1. ACCURACY COMPARISON:

- Custom CNN outperformed AlexNet by 10.41%
- This suggests that the custom architecture is better suited for CIFAR-10

2. MODEL EFFICIENCY:

- Custom CNN Efficiency: 0.582 (accuracy per million parameters)
- AlexNet Efficiency: 0.020 (accuracy per million parameters)
- Custom CNN is more parameter-efficient

3. ARCHITECTURE INSIGHTS:

- Custom CNN uses modern techniques: Batch Normalization, Dropout
- AlexNet adapted from 224x224 to 32x32 input size
- Both models benefit from data normalization and proper regularization

4. PRACTICAL CONSIDERATIONS:

- Custom CNN trains 3.5x faster
- Model size affects deployment and inference speed
- Custom architectures can be optimized for specific datasets

CONCLUSION: The Custom CNN model provides the best performance for CIFAR-10 classification in terms of test accuracy, achieving 85.74% accuracy.

=====