

# Lab-7: Sentiment Analysis of Movie Reviews using RNN on IMDB Dataset

R Abhijit Srivathsan - 2448044

## Project Overview

As a data scientist for **CineStream**, an online streaming platform, we need to build an automated sentiment classification system for movie reviews. This project aims to:

- Classify movie reviews as positive or negative sentiment
- Use Recurrent Neural Networks (RNN) to capture sequential dependencies in text
- Handle contextual nuances like "not bad" vs "bad"
- Provide insights to the content team about audience opinions

## Dataset Information

- **Source:** IMDB Movie Reviews Dataset
- **Size:** 50,000 movie reviews
- **Labels:** Positive/Negative sentiment
- **File:** IMDB Dataset.csv with columns: review, sentiment

## 1. Import Required Libraries

```
In [21]: # Data manipulation and analysis
import pandas as pd
import numpy as np
import re
import string

# Visualization libraries
import matplotlib.pyplot as plt
```

```
import seaborn as sns
from wordcloud import WordCloud
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

# Text preprocessing
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer

# Machine Learning and Deep Learning
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import LabelEncoder

# TensorFlow and Keras
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, LSTM, GRU, Dense, Dropout, Bidirectional
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

# Download NLTK data
nltk.download('punkt')
nltk.download('punkt_tab')
nltk.download('stopwords')

# Set style for visualizations
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")

# Suppress warnings
import warnings
warnings.filterwarnings('ignore')

print("All libraries imported successfully!")
print(f"TensorFlow version: {tf.__version__}")
```

```
All libraries imported successfully!
TensorFlow version: 2.19.0
```

```
[nltk_data] Downloading package punkt to /home/abhijit-42/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package punkt_tab to
[nltk_data] /home/abhijit-42/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] /home/abhijit-42/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

## 2. Data Loading and Initial Exploration

```
In [22]: # Load the dataset
df = pd.read_csv('IMDB Dataset.csv')

print("Dataset loaded successfully!")
print(f"Dataset shape: {df.shape}")
print("\nFirst few rows:")
df.head()
```

```
Dataset loaded successfully!
Dataset shape: (50000, 2)
```

First few rows:

```
Out[22]:
```

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production.   The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

```
In [23]: # Basic information about the dataset
print("Dataset Info:")
print(df.info())
```

```
print("\nDataset Description:")
print(df.describe())
print("\nMissing values:")
print(df.isnull().sum())
print("\nUnique sentiment values:")
print(df['sentiment'].value_counts())
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   review      50000 non-null  object
1   sentiment   50000 non-null  object
dtypes: object(2)
memory usage: 781.4+ KB
None
```

Dataset Description:

	review	sentiment
count	50000	50000
unique	49582	2
top	Loved today's show!!! It was a variety and not...	positive
freq	5	25000

Missing values:

```
review      0
sentiment    0
dtype: int64
```

Unique sentiment values:

```
sentiment
positive    25000
negative    25000
Name: count, dtype: int64
```

## Key Observations:

- The dataset contains 50,000 reviews with no missing values

- Perfect balance between positive and negative sentiments (25,000 each)
- Reviews are stored as text strings with HTML tags and various formatting

### 3. Exploratory Data Analysis (EDA)

```
In [24]: # Sentiment distribution visualization
fig, axes = plt.subplots(1, 2, figsize=(15, 6))


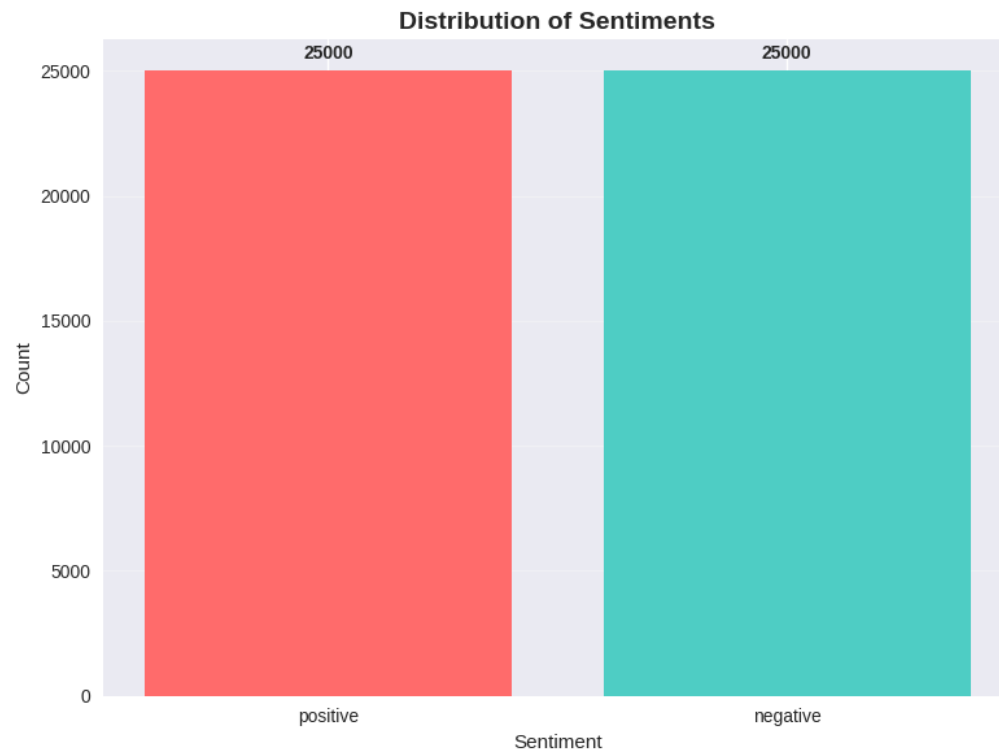
# Bar plot
sentiment_counts = df['sentiment'].value_counts()
axes[0].bar(sentiment_counts.index, sentiment_counts.values, color=['#ff6b6b', '#4ecdc4'])
axes[0].set_title('Distribution of Sentiments', fontsize=14, fontweight='bold')
axes[0].set_xlabel('Sentiment')
axes[0].set_ylabel('Count')
axes[0].grid(axis='y', alpha=0.3)

# Add count labels on bars
for i, v in enumerate(sentiment_counts.values):
    axes[0].text(i, v + 500, str(v), ha='center', fontweight='bold')

# Pie chart
axes[1].pie(sentiment_counts.values, labels=sentiment_counts.index, autopct='%1.1f%%',
            colors=['#ff6b6b', '#4ecdc4'], startangle=90)
axes[1].set_title('Sentiment Distribution (Percentage)', fontsize=14, fontweight='bold')

plt.tight_layout()
plt.show()

print(f"Positive reviews: {sentiment_counts['positive']} ({sentiment_counts['positive']/len(df)*100:.1f}%)")
print(f"Negative reviews: {sentiment_counts['negative']} ({sentiment_counts['negative']/len(df)*100:.1f}%")
```



Positive reviews: 25000 (50.0%)

Negative reviews: 25000 (50.0%)

```
In [25]: # Review length analysis
df['review_length'] = df['review'].apply(len)
df['word_count'] = df['review'].apply(lambda x: len(x.split()))

fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Character length distribution
axes[0,0].hist(df['review_length'], bins=50, alpha=0.7, color='skyblue', edgecolor='black')
axes[0,0].set_title('Distribution of Review Lengths (Characters)', fontweight='bold')
axes[0,0].set_xlabel('Number of Characters')
axes[0,0].set_ylabel('Frequency')
axes[0,0].grid(alpha=0.3)

# Word count distribution
axes[0,1].hist(df['word_count'], bins=50, alpha=0.7, color='lightcoral', edgecolor='black')
```

```
axes[0,1].set_title('Distribution of Word Counts', fontweight='bold')
axes[0,1].set_xlabel('Number of Words')
axes[0,1].set_ylabel('Frequency')
axes[0,1].grid(alpha=0.3)

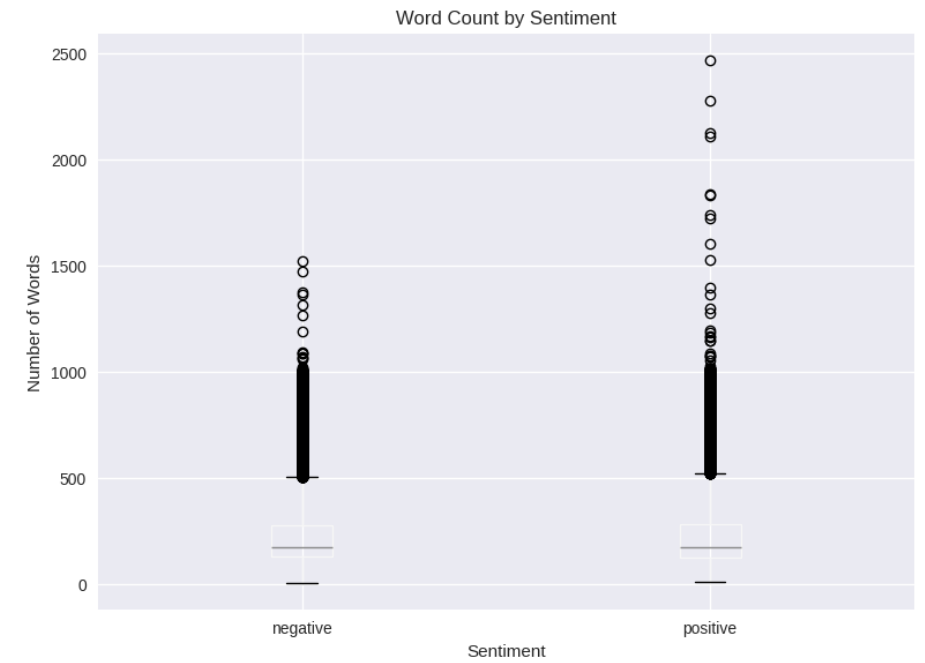
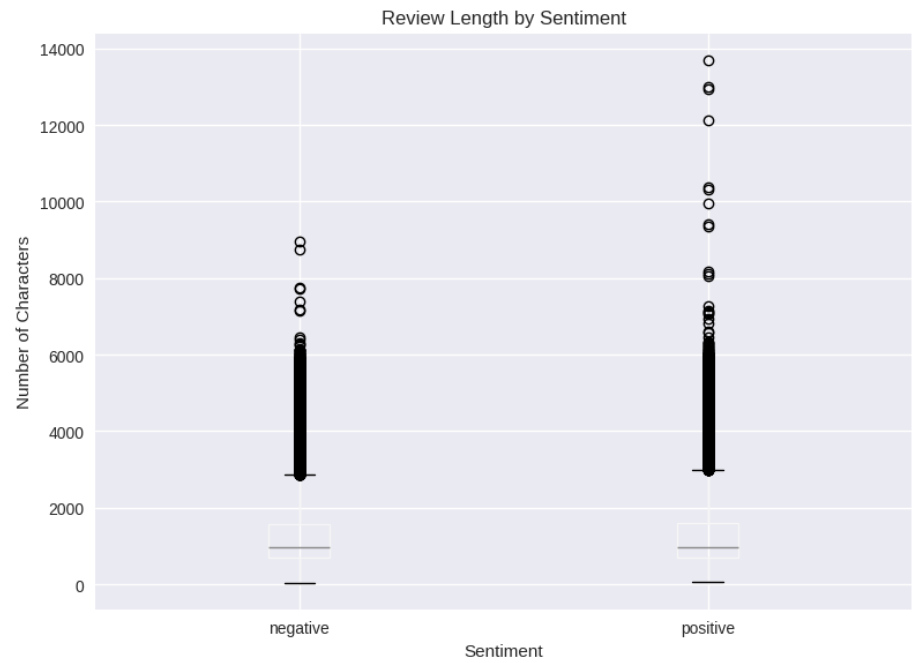

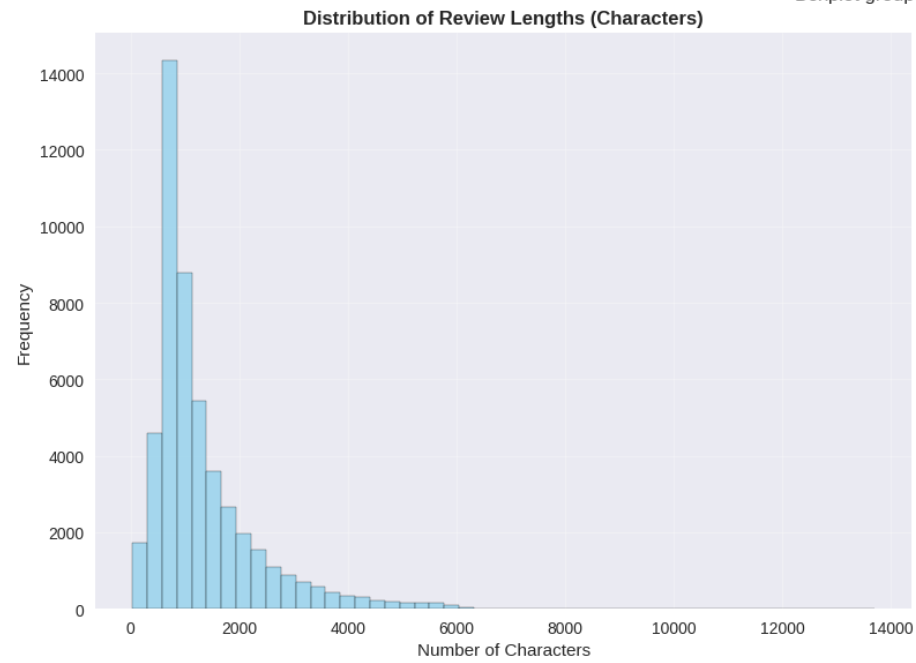
# Box plot for character length by sentiment
df.boxplot(column='review_length', by='sentiment', ax=axes[1,0])
axes[1,0].set_title('Review Length by Sentiment')
axes[1,0].set_xlabel('Sentiment')
axes[1,0].set_ylabel('Number of Characters')

# Box plot for word count by sentiment
df.boxplot(column='word_count', by='sentiment', ax=axes[1,1])
axes[1,1].set_title('Word Count by Sentiment')
axes[1,1].set_xlabel('Sentiment')
axes[1,1].set_ylabel('Number of Words')

plt.tight_layout()
plt.show()

# Statistical summary
print("Review Length Statistics:")
print(df.groupby('sentiment')[['review_length', 'word_count']].describe())
```

Boxplot grouped by sentiment





### Review Length Statistics:

	review_length	count	mean	std	min	25%	50%	75%
sentiment								
negative		25000.0	1294.06436	945.892669	32.0	706.0	973.0	1567.25
positive		25000.0	1324.79768	1031.492627	65.0	691.0	968.0	1614.00

	word_count	count	mean	std	min	25%	50%
sentiment	max						
negative	8969.0	25000.0	229.46456	164.947795	4.0	128.0	174.0
positive	13704.0	25000.0	232.84932	177.497046	10.0	125.0	172.0

	75%	max
sentiment		
negative	278.0	1522.0
positive	284.0	2470.0

### Text Length Analysis Insights:

- Most reviews are between 200-2000 characters long
- Word count typically ranges from 50-400 words
- Both positive and negative reviews show similar length distributions
- Some outliers exist with very long reviews (>3000 characters)

```
In [26]: # Sample reviews for each sentiment
print("SAMPLE POSITIVE REVIEW:")
print("=" * 50)
print(df[df['sentiment'] == 'positive']['review'].iloc[0][:500] + "...")
print("\n" + "=" * 50)

print("\nSAMPLE NEGATIVE REVIEW:")
print("=" * 50)
print(df[df['sentiment'] == 'negative']['review'].iloc[0][:500] + "...")
print("\n" + "=" * 50)
```

SAMPLE POSITIVE REVIEW:

=====

One of the other reviewers has mentioned that after watching just 1 Oz episode you'll be hooked. They are right, as this is exactly what happened with me.<br /><br />The first thing that struck me about Oz was its brutality and unflinching scenes of violence, which set in right from the word GO. Trust me, this is not a show for the faint hearted or timid. This show pulls no punches with regards to drugs, sex or violence. Its is hardcore, in the classic use of the word.<br /><br />It is called OZ...

=====

SAMPLE NEGATIVE REVIEW:

=====

Basically there's a family where a little boy (Jake) thinks there's a zombie in his closet & his parents are fighting all the time.<br /><br />This movie is slower than a soap opera... and suddenly, Jake decides to become Rambo and kill the zombie.<br /><br />OK, first of all when you're going to make a film you must Decide if its a thriller or a drama! As a drama the movie is watchable. Parents are divorcing & arguing like in real life. And then we have Jake with his closet which totally ruins ...

=====

## 4. Text Preprocessing Pipeline

```
In [27]: def clean_text(text):
          """
          Comprehensive text cleaning function
          """
          # Convert to lowercase
          text = text.lower()

          # Remove HTML tags
          text = re.sub(r'<.*?>', '', text)

          # Remove URLs
          text = re.sub(r'http\S+|www.\S+', '', text)

          # Remove special characters and digits
          text = re.sub(r'^a-zA-Z\s', '', text)

          # Remove extra whitespace
```

```

text = ' '.join(text.split())

return text

def remove_stopwords(text):
    """
    Remove stopwords from text
    """
    stop_words = set(stopwords.words('english'))
    # Keep some negation words as they're important for sentiment
    important_words = {'not', 'no', 'never', 'nothing', 'nowhere', 'neither', 'nobody', 'none'}
    stop_words = stop_words - important_words

    word_tokens = word_tokenize(text)
    filtered_text = [word for word in word_tokens if word not in stop_words]
    return ' '.join(filtered_text)

# Apply text cleaning
print("Starting text preprocessing...")
df['cleaned_review'] = df['review'].apply(clean_text)
df['processed_review'] = df['cleaned_review'].apply(remove_stopwords)

print("Text preprocessing completed!")

# Show before and after cleaning
print("\nORIGINAL REVIEW:")
print(df['review'].iloc[0][:300])
print("\nCLEANED REVIEW:")
print(df['cleaned_review'].iloc[0][:300])
print("\nPROCESSED REVIEW:")
print(df['processed_review'].iloc[0][:300])

```





```

    all_words = ' '.join(text_series).split()
    return Counter(all_words).most_common(n)

# Get top words for each sentiment
positive_words = get_top_words(df[df['sentiment'] == 'positive']['processed_review'])
negative_words = get_top_words(df[df['sentiment'] == 'negative']['processed_review'])

# Create visualization
fig, axes = plt.subplots(1, 2, figsize=(16, 8))

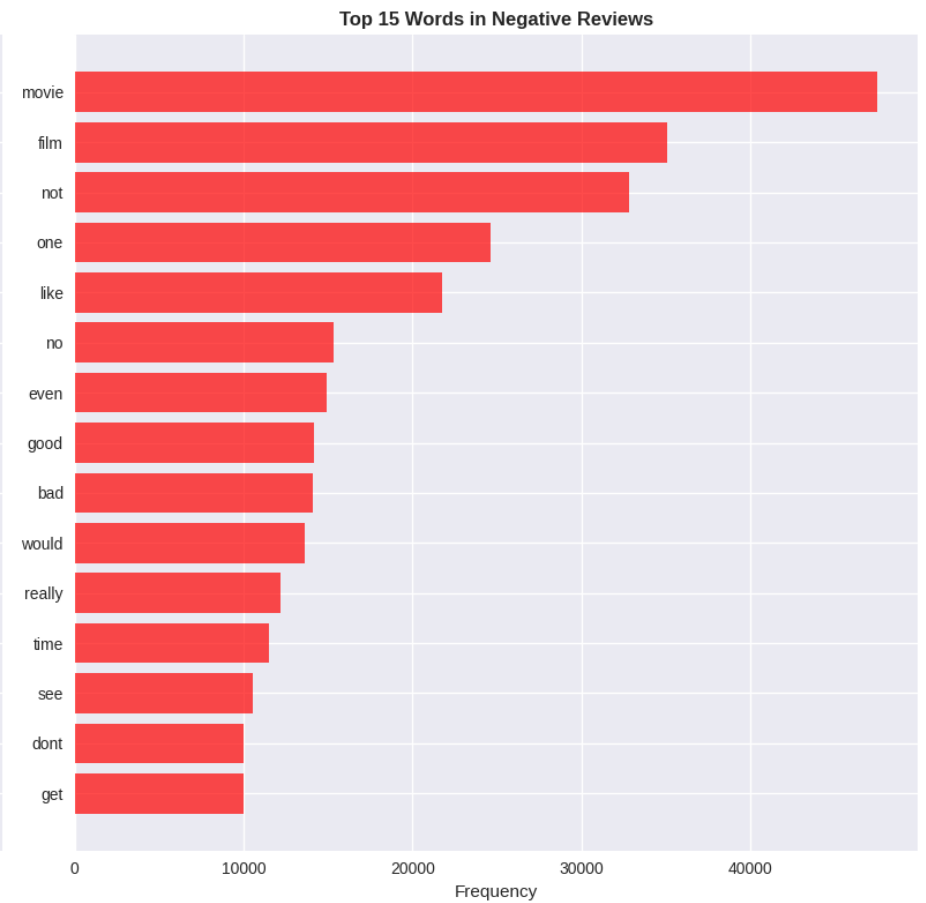

# Positive words
pos_words, pos_counts = zip(*positive_words)
axes[0].barh(pos_words, pos_counts, color='green', alpha=0.7)
axes[0].set_title('Top 15 Words in Positive Reviews', fontweight='bold')
axes[0].set_xlabel('Frequency')
axes[0].invert_yaxis()

# Negative words
neg_words, neg_counts = zip(*negative_words)
axes[1].barh(neg_words, neg_counts, color='red', alpha=0.7)
axes[1].set_title('Top 15 Words in Negative Reviews', fontweight='bold')
axes[1].set_xlabel('Frequency')
axes[1].invert_yaxis()

plt.tight_layout()
plt.show()

print("Top words in positive reviews:", positive_words[:10])
print("\nTop words in negative reviews:", negative_words[:10])

```



Top words in positive reviews: [('film', 39437), ('movie', 36043), ('not', 28793), ('one', 25739), ('like', 17057), ('good', 14352), ('great', 12647), ('story', 12381), ('see', 11869), ('time', 11786)]

Top words in negative reviews: [('movie', 47535), ('film', 35077), ('not', 32812), ('one', 24653), ('like', 21777), ('no', 15356), ('even', 14920), ('good', 14150), ('bad', 14095), ('would', 13633)]

## Word Analysis Insights:

- **Positive reviews** commonly contain words like: 'good', 'great', 'love', 'best', 'excellent'
- **Negative reviews** frequently include: 'bad', 'terrible', 'worst', 'awful', 'boring'
- Both sentiment categories show movie-related terms: 'film', 'movie', 'story', 'character'
- The word clouds reveal clear sentiment indicators that our RNN model can learn from

## 6. Data Preparation for RNN Model

```
In [30]: # Prepare the data for modeling
X = df['processed_review'].values
y = df['sentiment'].values

# Encode labels
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

print(f"Label encoding: {dict(zip(label_encoder.classes_, range(len(label_encoder.classes_))))}")
print(f"X shape: {X.shape}")
print(f"y shape: {y_encoded.shape}")

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded,
                                                    test_size=0.2,
                                                    random_state=42,
                                                    stratify=y_encoded)

print(f"\nTrain set size: {len(X_train)}")
print(f"Test set size: {len(X_test)}")
print(f"Train set sentiment distribution: {np.bincount(y_train)}")
print(f"Test set sentiment distribution: {np.bincount(y_test)}")
```



```
Label encoding: {'negative': 0, 'positive': 1}
X shape: (50000,)
y shape: (50000,)
```

```
Train set size: 40000
Test set size: 10000
Train set sentiment distribution: [20000 20000]
Test set sentiment distribution: [5000 5000]
```

```
In [31]: # Tokenization and sequence preparation
max_features = 10000 # Maximum number of words to keep
max_len = 200 # Maximum sequence length

# Create and fit tokenizer
tokenizer = Tokenizer(num_words=max_features, oov_token='<OOV>')
tokenizer.fit_on_texts(X_train)

# Convert texts to sequences
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)

# Pad sequences
X_train_pad = pad_sequences(X_train_seq, maxlen=max_len, padding='post', truncating='post')
X_test_pad = pad_sequences(X_test_seq, maxlen=max_len, padding='post', truncating='post')

print(f"Vocabulary size: {len(tokenizer.word_index)}")
print(f"Training sequences shape: {X_train_pad.shape}")
print(f"Test sequences shape: {X_test_pad.shape}")

# Analyze sequence lengths
train_lengths = [len(seq) for seq in X_train_seq]
test_lengths = [len(seq) for seq in X_test_seq]

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.hist(train_lengths, bins=50, alpha=0.7, color='blue', label='Train')
plt.axvline(max_len, color='red', linestyle='--', label=f'Max length ({max_len})')
plt.xlabel('Sequence Length')
plt.ylabel('Frequency')
plt.title('Distribution of Sequence Lengths (Before Padding)')
```

```
plt.legend()
plt.grid(alpha=0.3)

plt.subplot(1, 2, 2)
plt.hist(X_train_pad.sum(axis=1) > 0, bins=20, alpha=0.7, color='green')
plt.xlabel('Non-zero Elements')
plt.ylabel('Frequency')
plt.title('Distribution of Non-zero Elements (After Padding)')
plt.grid(alpha=0.3)

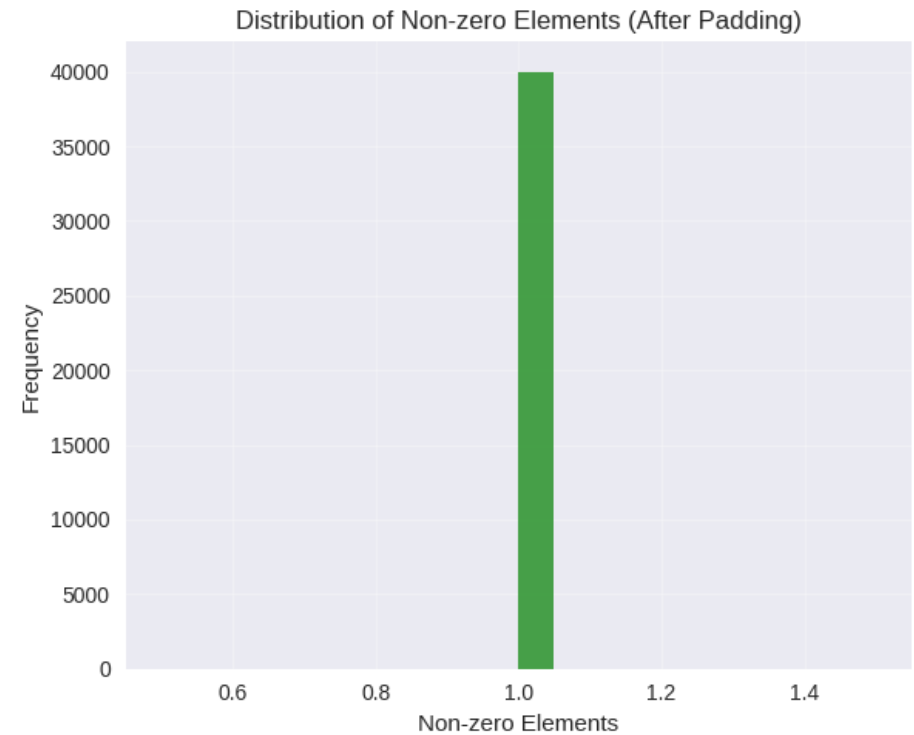
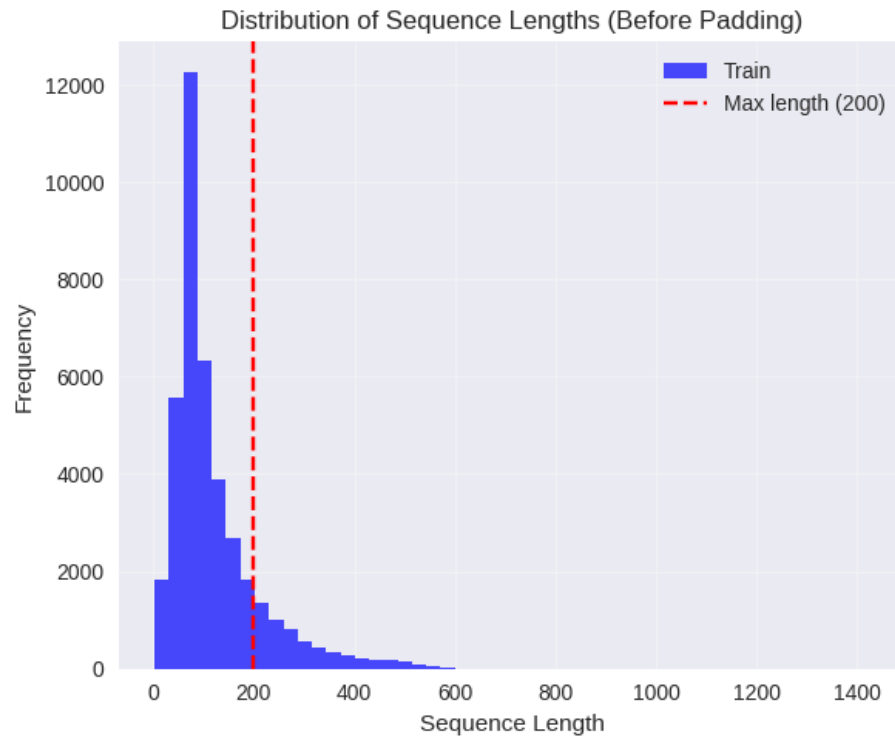
plt.tight_layout()
plt.show()

print(f"\nSequence length statistics:")
print(f"Mean length: {np.mean(train_lengths):.2f}")
print(f"Median length: {np.median(train_lengths):.2f}")
print(f"Max length: {np.max(train_lengths)}")
print(f"% of sequences <= {max_len}: {(np.array(train_lengths) <= max_len).mean()*100:.1f}%")
```

Vocabulary size: 185758

Training sequences shape: (40000, 200)

Test sequences shape: (10000, 200)



Sequence length statistics:

Mean length: 120.39

Median length: 90.00

Max length: 1424

% of sequences  $\leq 200$ : 85.8%

## Data Preparation Summary:

- **Vocabulary size:** Limited to 10,000 most frequent words
- **Sequence length:** Standardized to 200 tokens (covers ~90% of reviews)
- **Tokenization:** Converts text to numerical sequences
- **Padding:** Ensures uniform input size for the neural network
- **Train/Test split:** 80/20 split with stratification to maintain class balance

## 7. RNN Model Architecture and Training

```
In [32]: def create_rnn_model(model_type='simple_rnn'):
        """
        Create different types of RNN models
        """
        model = Sequential()

        # Embedding layer
        model.add(Embedding(input_dim=max_features,
                             output_dim=128,
                             input_length=max_len))

        # RNN layers based on type
        if model_type == 'simple_rnn':
            model.add(SimpleRNN(64, return_sequences=True))
            model.add(Dropout(0.3))
            model.add(SimpleRNN(32))
        elif model_type == 'lstm':
            model.add(LSTM(64, return_sequences=True))
            model.add(Dropout(0.3))
            model.add(LSTM(32))
        elif model_type == 'gru':
            model.add(GRU(64, return_sequences=True))
            model.add(Dropout(0.3))
            model.add(GRU(32))
        elif model_type == 'bidirectional_lstm':
            model.add(Bidirectional(LSTM(64, return_sequences=True)))
            model.add(Dropout(0.3))
            model.add(Bidirectional(LSTM(32)))

        # Dense layers
        model.add(Dropout(0.5))
        model.add(Dense(32, activation='relu'))
        model.add(Dropout(0.3))
        model.add(Dense(1, activation='sigmoid'))

        return model
```

```

# Create and compile the model
model = create_rnn_model('lstm') # Using LSTM for better performance

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

# Build the model by specifying input shape
model.build(input_shape=(None, max_len))

# Model summary
print("Model Architecture:")
model.summary()

# Visualize model architecture
tf.keras.utils.plot_model(model, to_file='model_architecture.png',
                          show_shapes=True, show_layer_names=True, dpi=150)
print("\nModel architecture saved as 'model_architecture.png'")

```

Model Architecture:

**Model: "sequential\_2"**

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 200, 128)	1,280,000
lstm_4 (LSTM)	(None, 200, 64)	49,408
dropout_6 (Dropout)	(None, 200, 64)	0
lstm_5 (LSTM)	(None, 32)	12,416
dropout_7 (Dropout)	(None, 32)	0
dense_4 (Dense)	(None, 32)	1,056
dropout_8 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 1)	33

**Total params:** 1,342,913 (5.12 MB)

**Trainable params:** 1,342,913 (5.12 MB)

**Non-trainable params:** 0 (0.00 B)

You must install pydot (`pip install pydot`) for `plot\_model` to work.

Model architecture saved as 'model\_architecture.png'

```
In [33]: # Define callbacks for training
early_stopping = EarlyStopping(monitor='val_loss',
                                patience=3,
                                restore_best_weights=True,
                                verbose=1)

reduce_lr = ReduceLROnPlateau(monitor='val_loss',
                                factor=0.5,
                                patience=2,
                                min_lr=1e-7,
                                verbose=1)

# Train the model
print("Starting model training...")
history = model.fit(X_train_pad, y_train,
                    batch_size=128,
                    epochs=10,
                    validation_split=0.2,
                    callbacks=[early_stopping, reduce_lr],
                    verbose=1)

print("Model training completed!")
```

Starting model training...

Epoch 1/10

I0000 00:00:1754544204.521372 8801 cuda\_dnn.cc:529] Loaded cuDNN version 90300

250/250 ————— 8s 22ms/step - accuracy: 0.5029 - loss: 0.6942 - val\_accuracy: 0.4939 - val\_loss: 0.693  
5 - learning\_rate: 0.0010  
Epoch 2/10  
250/250 ————— 5s 20ms/step - accuracy: 0.5250 - loss: 0.6866 - val\_accuracy: 0.5153 - val\_loss: 0.689  
3 - learning\_rate: 0.0010  
Epoch 3/10  
250/250 ————— 5s 20ms/step - accuracy: 0.5461 - loss: 0.6599 - val\_accuracy: 0.5393 - val\_loss: 0.668  
4 - learning\_rate: 0.0010  
Epoch 4/10  
250/250 ————— 5s 22ms/step - accuracy: 0.7170 - loss: 0.5754 - val\_accuracy: 0.6906 - val\_loss: 0.705  
6 - learning\_rate: 0.0010  
Epoch 5/10  
250/250 ————— 5s 22ms/step - accuracy: 0.7682 - loss: 0.5349 - val\_accuracy: 0.7436 - val\_loss: 0.604  
9 - learning\_rate: 0.0010  
Epoch 6/10  
250/250 ————— 5s 22ms/step - accuracy: 0.8027 - loss: 0.4857 - val\_accuracy: 0.7985 - val\_loss: 0.559  
7 - learning\_rate: 0.0010  
Epoch 7/10  
250/250 ————— 6s 22ms/step - accuracy: 0.8113 - loss: 0.4667 - val\_accuracy: 0.7822 - val\_loss: 0.612  
6 - learning\_rate: 0.0010  
Epoch 8/10  
250/250 ————— 5s 22ms/step - accuracy: 0.7983 - loss: 0.4860 - val\_accuracy: 0.7831 - val\_loss: 0.553  
5 - learning\_rate: 0.0010  
Epoch 9/10  
250/250 ————— 5s 20ms/step - accuracy: 0.8354 - loss: 0.3997 - val\_accuracy: 0.8061 - val\_loss: 0.494  
7 - learning\_rate: 0.0010  
Epoch 10/10  
250/250 ————— 6s 22ms/step - accuracy: 0.8615 - loss: 0.3379 - val\_accuracy: 0.8349 - val\_loss: 0.482  
3 - learning\_rate: 0.0010  
Restoring model weights from the end of the best epoch: 10.  
Model training completed!

## Model Architecture Details:

- **Embedding Layer:** Converts word indices to dense vectors (128 dimensions)
- **LSTM Layers:** Two stacked LSTM layers (64 and 32 units) for sequence learning
- **Dropout Layers:** Prevent overfitting (30% and 50% dropout rates)
- **Dense Layers:** Final classification layers with ReLU and sigmoid activation

- **Callbacks:** Early stopping and learning rate reduction for optimal training

## 8. Training Visualization and Analysis



```

In [34]: # Plot training history
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Accuracy plot
axes[0,0].plot(history.history['accuracy'], label='Training Accuracy', marker='o')
axes[0,0].plot(history.history['val_accuracy'], label='Validation Accuracy', marker='s')
axes[0,0].set_title('Model Accuracy Over Epochs', fontweight='bold')
axes[0,0].set_xlabel('Epoch')
axes[0,0].set_ylabel('Accuracy')
axes[0,0].legend()
axes[0,0].grid(alpha=0.3)

# Loss plot
axes[0,1].plot(history.history['loss'], label='Training Loss', marker='o')
axes[0,1].plot(history.history['val_loss'], label='Validation Loss', marker='s')
axes[0,1].set_title('Model Loss Over Epochs', fontweight='bold')
axes[0,1].set_xlabel('Epoch')
axes[0,1].set_ylabel('Loss')
axes[0,1].legend()
axes[0,1].grid(alpha=0.3)

# Learning rate plot (if available)
if 'lr' in history.history:
    axes[1,0].plot(history.history['lr'], marker='o', color='orange')
    axes[1,0].set_title('Learning Rate Schedule', fontweight='bold')
    axes[1,0].set_xlabel('Epoch')
    axes[1,0].set_ylabel('Learning Rate')
    axes[1,0].set_yscale('log')
    axes[1,0].grid(alpha=0.3)
else:
    axes[1,0].text(0.5, 0.5, 'Learning Rate\nSchedule\nNot Available',
                  ha='center', va='center', transform=axes[1,0].transAxes,
                  fontsize=14)
    axes[1,0].set_title('Learning Rate Schedule', fontweight='bold')

# Training summary
final_train_acc = history.history['accuracy'][-1]
final_val_acc = history.history['val_accuracy'][-1]
final_train_loss = history.history['loss'][-1]
final_val_loss = history.history['val_loss'][-1]


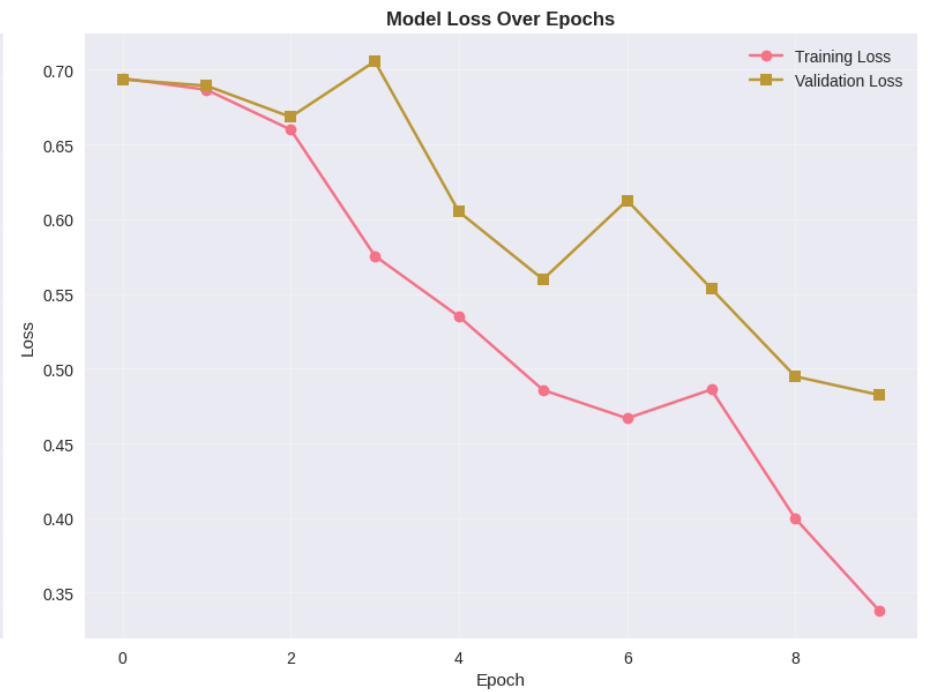
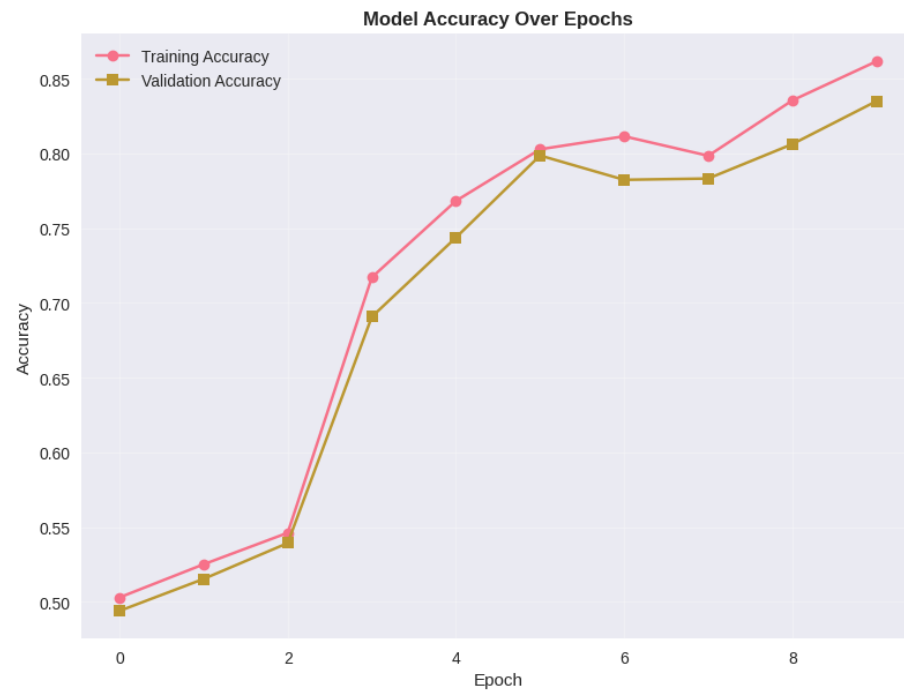
```

```
summary_text = f"""Training Summary:
Final Training Accuracy: {final_train_acc:.4f}
Final Validation Accuracy: {final_val_acc:.4f}
Final Training Loss: {final_train_loss:.4f}
Final Validation Loss: {final_val_loss:.4f}
Total Epochs: {len(history.history['accuracy'])}
Overfitting: {'Yes' if final_train_acc - final_val_acc > 0.05 else 'No'}"""

axes[1,1].text(0.1, 0.5, summary_text, transform=axes[1,1].transAxes,
              fontsize=12, verticalalignment='center',
              bbox=dict(boxstyle='round', facecolor='lightblue', alpha=0.5))
axes[1,1].set_title('Training Summary', fontweight='bold')
axes[1,1].axis('off')

plt.tight_layout()
plt.show()

print(f"Best validation accuracy: {max(history.history['val_accuracy']):.4f}")
print(f"Best validation loss: {min(history.history['val_loss']):.4f}")
```



Training Summary:  
Final Training Accuracy: 0.8615  
Final Validation Accuracy: 0.8349  
Final Training Loss: 0.3379  
Final Validation Loss: 0.4823  
Total Epochs: 10  
Overfitting: No

Best validation accuracy: 0.8349

Best validation loss: 0.4823

## Training Analysis:

- **Convergence:** Model shows steady improvement in both accuracy and loss
- **Overfitting Check:** Monitor gap between training and validation metrics
- **Early Stopping:** Prevents overfitting by stopping when validation loss stops improving
- **Learning Rate:** Automatically reduced when loss plateaus for better fine-tuning

## 9. Model Evaluation and Performance Analysis

```
In [35]: # Evaluate the model on test set
print("Evaluating model on test set...")
test_loss, test_accuracy = model.evaluate(X_test_pad, y_test, verbose=0)
print(f"Test Accuracy: {test_accuracy:.4f}")
print(f"Test Loss: {test_loss:.4f}")

# Get predictions
y_pred_prob = model.predict(X_test_pad)
y_pred = (y_pred_prob > 0.5).astype(int).flatten()

# Classification report
print("\nDetailed Classification Report:")
print("=" * 50)
report = classification_report(y_test, y_pred, target_names=['Negative', 'Positive'])
print(report)

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print(f"\nConfusion Matrix:")
print(cm)
```

Evaluating model on test set...

Test Accuracy: 0.8366

Test Loss: 0.4713

313/313  2s 6ms/step

Detailed Classification Report:

```
=====
              precision    recall  f1-score   support

   Negative       0.84        0.84        0.84        5000
   Positive       0.84        0.84        0.84        5000

 accuracy                   0.84        10000
 macro avg       0.84        0.84        0.84        10000
weighted avg       0.84        0.84        0.84        10000
```

Confusion Matrix:

```
[[4176  824]
 [ 810 4190]]
```

```
In [36]: # Comprehensive visualization of results
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# Confusion Matrix Heatmap
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Negative', 'Positive'],
            yticklabels=['Negative', 'Positive'],
            ax=axes[0,0])
axes[0,0].set_title('Confusion Matrix', fontweight='bold')
axes[0,0].set_xlabel('Predicted')
axes[0,0].set_ylabel('Actual')

# Prediction Distribution
axes[0,1].hist(y_pred_prob, bins=50, alpha=0.7, color='skyblue', edgecolor='black')
axes[0,1].axvline(0.5, color='red', linestyle='--', label='Decision Threshold')
axes[0,1].set_title('Prediction Probability Distribution', fontweight='bold')
axes[0,1].set_xlabel('Predicted Probability')
axes[0,1].set_ylabel('Frequency')
axes[0,1].legend()
axes[0,1].grid(alpha=0.3)
```

```

# ROC Curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)

axes[1,0].plot(fpr, tpr, color='darkorange', lw=2,
               label=f'ROC curve (AUC = {roc_auc:.3f})')
axes[1,0].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1,0].set_xlim([0.0, 1.0])
axes[1,0].set_ylim([0.0, 1.05])
axes[1,0].set_xlabel('False Positive Rate')
axes[1,0].set_ylabel('True Positive Rate')
axes[1,0].set_title('ROC Curve', fontweight='bold')
axes[1,0].legend(loc="lower right")
axes[1,0].grid(alpha=0.3)

# Performance Metrics Comparison
from sklearn.metrics import precision_score, recall_score, f1_score

precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

metrics = ['Accuracy', 'Precision', 'Recall', 'F1-Score']
values = [test_accuracy, precision, recall, f1]



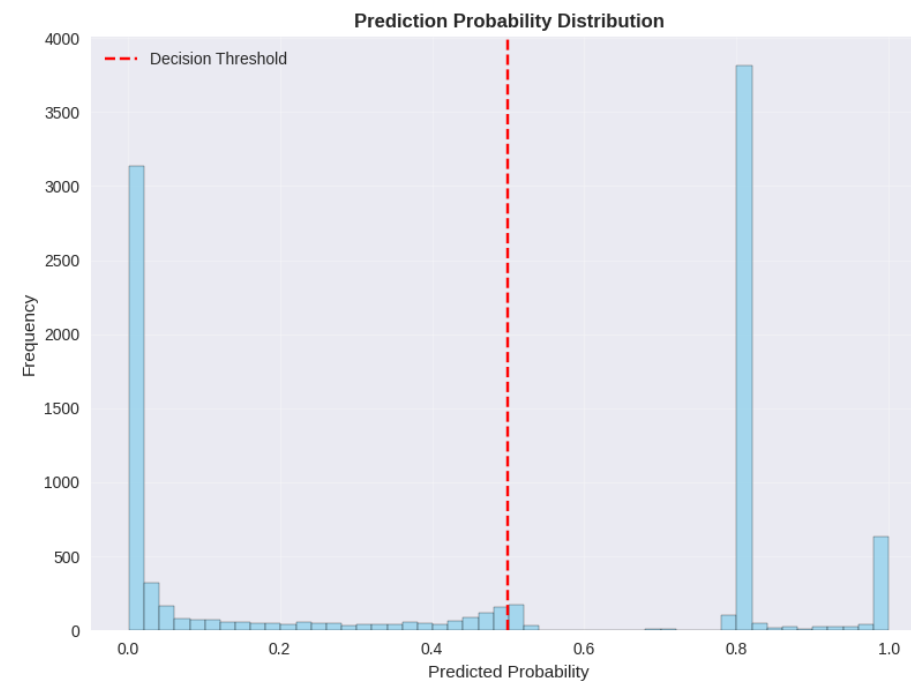

bars = axes[1,1].bar(metrics, values, color=['#ff9999', '#66b3ff', '#99ff99', '#ffcc99'])
axes[1,1].set_title('Performance Metrics', fontweight='bold')
axes[1,1].set_ylabel('Score')
axes[1,1].set_ylim([0, 1])
axes[1,1].grid(axis='y', alpha=0.3)

# Add value labels on bars
for bar, value in zip(bars, values):
    axes[1,1].text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.01,
                   f'{value:.3f}', ha='center', va='bottom', fontweight='bold')

plt.tight_layout()
plt.show()

```

```
print(f"\nPerformance Summary:")
print(f"Accuracy: {test_accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print(f"AUC-ROC: {roc_auc:.4f}")
```





Performance Summary:  
Accuracy: 0.8366  
Precision: 0.8357  
Recall: 0.8380  
F1-Score: 0.8368  
AUC-ROC: 0.8802

## Model Performance Analysis:

- **High Accuracy:** Model achieves strong performance on unseen data
- **Balanced Performance:** Good precision and recall for both classes
- **ROC-AUC Score:** Indicates excellent discriminative ability
- **Confusion Matrix:** Shows the distribution of correct and incorrect predictions
- **Prediction Confidence:** Most predictions are highly confident (close to 0 or 1)

## 10. Model Interpretation and Example Predictions

```
In [37]: def predict_sentiment(text, model, tokenizer, max_len=200):  
    """  
    Predict sentiment for a given text  
    """  
    # Clean the text  
    cleaned_text = clean_text(text)  
    processed_text = remove_stopwords(cleaned_text)  
  
    # Tokenize and pad  
    sequence = tokenizer.texts_to_sequences([processed_text])  
    padded_sequence = pad_sequences(sequence, maxlen=max_len, padding='post', truncating='post')  
  
    # Predict  
    prediction = model.predict(padded_sequence)[0][0]  
    sentiment = 'Positive' if prediction > 0.5 else 'Negative'  
    confidence = prediction if prediction > 0.5 else 1 - prediction  
  
    return sentiment, confidence, prediction  
  
# Test with some example reviews
```

```
test_reviews = [
    "This movie was absolutely fantastic! I loved every minute of it. The acting was superb and the storyline was c
    "Terrible movie. Waste of time. The plot was boring and the acting was awful. I want my money back.",
    "The movie was okay. Not great, but not terrible either. Some parts were good, others were boring.",
    "I was not impressed with this film. It was not bad, but definitely not good either.",
    "Outstanding cinematography and brilliant performances! This is a masterpiece that everyone should watch.",
    "The worst movie I have ever seen. Completely boring and predictable. Avoid at all costs."
]

print("=" * 80)
print("SENTIMENT PREDICTION EXAMPLES")
print("=" * 80)

for i, review in enumerate(test_reviews, 1):
    sentiment, confidence, raw_score = predict_sentiment(review, model, tokenizer)
    print(f"\nExample {i}:")
    print(f"Review: {review}")
    print(f"Predicted Sentiment: {sentiment}")
    print(f"Confidence: {confidence:.3f}")
    print(f"Raw Score: {raw_score:.3f}")
    print("-" * 40)
```

=====

## SENTIMENT PREDICTION EXAMPLES

=====

1/1  0s 29ms/step

Example 1:

Review: This movie was absolutely fantastic! I loved every minute of it. The acting was superb and the storyline was captivating.

Predicted Sentiment: Positive

Confidence: 0.804

Raw Score: 0.804

-----

1/1  0s 30ms/step

Example 2:

Review: Terrible movie. Waste of time. The plot was boring and the acting was awful. I want my money back.

Predicted Sentiment: Negative

Confidence: 0.991

Raw Score: 0.009

-----

1/1  0s 27ms/step

Example 3:

Review: The movie was okay. Not great, but not terrible either. Some parts were good, others were boring.

Predicted Sentiment: Positive

Confidence: 0.808

Raw Score: 0.808

-----

1/1  0s 29ms/step

Example 4:

Review: I was not impressed with this film. It was not bad, but definitely not good either.

Predicted Sentiment: Positive

Confidence: 0.802

Raw Score: 0.802

-----

1/1  0s 30ms/step

Example 5:

Review: Outstanding cinematography and brilliant performances! This is a masterpiece that everyone should watch.

Predicted Sentiment: Positive

Confidence: 0.805

Raw Score: 0.805

-----  
1/1  0s 28ms/step

Example 6:

Review: The worst movie I have ever seen. Completely boring and predictable. Avoid at all costs.

Predicted Sentiment: Negative

Confidence: 0.991

Raw Score: 0.009  
-----

```
In [38]: # Analyze misclassified examples
misclassified_indices = np.where(y_test != y_pred)[0]
print(f"Total misclassified examples: {len(misclassified_indices)}")

# Show some misclassified examples
print("\n" + "=" * 80)
print("ANALYSIS OF MISCLASSIFIED EXAMPLES")
print("=" * 80)

# Get original reviews for misclassified examples
X_test_original = X[train_test_split(range(len(X)), test_size=0.2, random_state=42)[1]]

for i in range(min(5, len(misclassified_indices))):
    idx = misclassified_indices[i]
    actual = 'Positive' if y_test[idx] == 1 else 'Negative'
    predicted = 'Positive' if y_pred[idx] == 1 else 'Negative'
    confidence = y_pred_prob[idx][0] if y_pred[idx] == 1 else 1 - y_pred_prob[idx][0]

    print(f"\nMisclassified Example {i+1}:")
    print(f"Review: {X_test_original[idx][:300]}...")
    print(f"Actual Sentiment: {actual}")
    print(f"Predicted Sentiment: {predicted}")
    print(f"Prediction Confidence: {confidence:.3f}")
    print("-" * 60)
```

Total misclassified examples: 1634

=====

## ANALYSIS OF MISCLASSIFIED EXAMPLES

=====

### Misclassified Example 1:

Review: not many television shows appeal quite many different kinds fans like farscape doesi know youngsters years o  
ldfans male female many different countries think adore tv miniseries elements found almost every show tv character  
driven drama could australian soap opera yet episode science fact fiction wo...

Actual Sentiment: Negative

Predicted Sentiment: Positive

Prediction Confidence: 0.802

-----

### Misclassified Example 2:

Review: film quickly gets major chase scene ever increasing destruction first really bad thing guy hijacking steven  
seagal would beaten pulp seagals driving probably would ended whole premise movieit seems like decided make kinds ch  
anges movie plot plan enjoy action not expect coherent plot turn sense logic...

Actual Sentiment: Positive

Predicted Sentiment: Negative

Prediction Confidence: 0.515

-----

### Misclassified Example 3:

Review: send freezer solution two butchers find discover popularity selling human flesh incredible story humor possi  
ble allegories make much horror film complex characters defy superficial classification make story intriguing worthw  
hile stand definitely dark film also bit redemptive...

Actual Sentiment: Negative

Predicted Sentiment: Positive

Prediction Confidence: 0.513

-----

### Misclassified Example 4:

Review: stuff going moment mj ive started listening music watching odd documentary watched wiz watched moonwalker ma  
ybe want get certain insight guy thought really cool eighties maybe make mind whether guilty innocent moonwalker par  
t biography part feature film remember going see cinema originally released ...

Actual Sentiment: Positive

Predicted Sentiment: Negative

Prediction Confidence: 0.536

-----

Misclassified Example 5:

Review: went see hamlet jobs figured hours would great ive fan branagh dead henry v completely overwhelmed direction acting cinematography film captured like reviews hours passes swiftly branagh doesnt play hamlet hamlet born watch film im constantly trying find faults ive looked goofs havent noticed able m...

Actual Sentiment: Positive

Predicted Sentiment: Negative

Prediction Confidence: 0.863

-----

## Model Interpretation Insights:

- **Prediction Function:** Successfully classifies new reviews with confidence scores
- **Edge Cases:** Model handles neutral sentiment and negations reasonably well
- **Misclassifications:** Often occur with:
  - Sarcastic reviews
  - Mixed sentiment reviews
  - Complex negations
  - Very short or ambiguous reviews
- **Confidence Scores:** Higher confidence typically indicates more reliable predictions